



Transact-SQL (TSQL) Migration Guide

Version 2019.4
2020-03-02

Transact-SQL (TSQL) Migration Guide

InterSystems IRIS Data Platform Version 2019.4 2020-03-02

Copyright © 2020 InterSystems Corporation

All rights reserved.

InterSystems, InterSystems IRIS, InterSystems Caché, InterSystems Ensemble, and InterSystems HealthShare are registered trademarks of InterSystems Corporation.

All other brand or product names used herein are trademarks or registered trademarks of their respective companies or organizations.

This document contains trade secret and confidential information which is the property of InterSystems Corporation, One Memorial Drive, Cambridge, MA 02142, or its affiliates, and is furnished for the sole purpose of the operation and maintenance of the products of InterSystems Corporation. No part of this publication is to be used for any other purpose, and this publication is not to be reproduced, copied, disclosed, transmitted, stored in a retrieval system or translated into any human or computer language, in any form, by any means, in whole or in part, without the express prior written consent of InterSystems Corporation.

The copying, use and disposition of this document and the software programs described herein is prohibited except to the limited extent set forth in the standard software license agreement(s) of InterSystems Corporation covering such programs and related documentation. InterSystems Corporation makes no representations and warranties concerning such software programs other than those set forth in such standard software license agreement(s). In addition, the liability of InterSystems Corporation for any losses or damages relating to or arising out of the use of such software programs is limited in the manner set forth in such standard software license agreement(s).

THE FOREGOING IS A GENERAL SUMMARY OF THE RESTRICTIONS AND LIMITATIONS IMPOSED BY INTERSYSTEMS CORPORATION ON THE USE OF, AND LIABILITY ARISING FROM, ITS COMPUTER SOFTWARE. FOR COMPLETE INFORMATION REFERENCE SHOULD BE MADE TO THE STANDARD SOFTWARE LICENSE AGREEMENT(S) OF INTERSYSTEMS CORPORATION, COPIES OF WHICH WILL BE MADE AVAILABLE UPON REQUEST.

InterSystems Corporation disclaims responsibility for errors which may appear in this document, and it reserves the right, in its sole discretion and without notice, to make substitutions and modifications in the products and practices described in this document.

For Support questions about any InterSystems products, contact:

InterSystems Worldwide Response Center (WRC)

Tel: +1-617-621-0700

Tel: +44 (0) 844 854 2917

Email: support@InterSystems.com

Table of Contents

About This Book	1
1 Overview	3
1.1 Migrating Existing TSQL Applications	3
1.1.1 Configuring TSQL	3
1.1.2 Migrating Source Code	3
1.1.3 Migrating the Data	4
1.2 Writing and Executing TSQL on InterSystems IRIS	4
2 InterSystems TSQL Constructs	7
2.1 Table References	7
2.2 Temporary Tables	7
2.3 System Tables	8
2.4 Transactions	8
2.5 Cursor Name Management	8
2.6 SYSOBJECTS References	9
3 InterSystems TSQL Language Elements	11
3.1 Literals	11
3.1.1 String Literals	11
3.1.2 Empty Strings	11
3.1.3 NULL	12
3.1.4 Hexadecimal	12
3.1.5 Reserved Words	12
3.1.6 Comments, Blank Lines, and Semicolons	12
3.2 Identifiers	13
3.2.1 Delimited and Quoted Identifiers	14
3.3 Data Types	14
3.4 Operators	15
3.4.1 Arithmetic and Equality Operators	15
3.4.2 Concatenation Operator	15
3.4.3 Comparison Operators	16
3.4.4 NOT Logical Operator	16
3.4.5 Bitwise Logical Operators	16
4 TSQL Commands	17
4.1 Data Definition Language (DDL) Statements	18
4.1.1 CREATE TABLE	18
4.1.2 ALTER TABLE	19
4.1.3 DROP TABLE	20
4.1.4 CREATE INDEX	21
4.1.5 DROP INDEX	21
4.1.6 CREATE TRIGGER	22
4.1.7 DROP TRIGGER	22
4.1.8 CREATE VIEW	22
4.1.9 DROP VIEW	23
4.1.10 CREATE DATABASE	23
4.1.11 DROP DATABASE	23
4.2 Data Management Language (DML) Statements	23

4.2.1 DELETE	24
4.2.2 INSERT	24
4.2.3 UPDATE	25
4.2.4 READTEXT	26
4.2.5 WRITETEXT	27
4.2.6 UPDATETEXT	27
4.2.7 TRUNCATE TABLE	28
4.3 Query Statements	28
4.3.1 SELECT	28
4.3.2 JOIN	31
4.3.3 UNION	31
4.3.4 FETCH Cursor	31
4.4 Flow of Control Statements	31
4.4.1 IF	31
4.4.2 WHILE	33
4.4.3 CASE	34
4.4.4 GOTO and Labels	34
4.4.5 WAITFOR	34
4.5 Assignment Statements	35
4.5.1 DECLARE	35
4.5.2 SET	35
4.6 Transaction Statements	36
4.6.1 SET TRANSACTION ISOLATION LEVEL	36
4.6.2 BEGIN TRANSACTION	36
4.6.3 COMMIT TRANSACTION	37
4.6.4 ROLLBACK TRANSACTION	37
4.6.5 SAVE TRANSACTION	38
4.6.6 LOCK TABLE	38
4.7 Procedure Statements	38
4.7.1 CREATE PROCEDURE / CREATE FUNCTION	38
4.7.2 ALTER FUNCTION	40
4.7.3 DROP FUNCTION	40
4.7.4 DROP PROCEDURE	40
4.7.5 RETURN	40
4.7.6 EXECUTE	40
4.7.7 CALL	41
4.8 Other Statements	41
4.8.1 CREATE USER	41
4.8.2 GRANT	42
4.8.3 REVOKE	42
4.8.4 PRINT	42
4.8.5 RAISERROR	43
4.8.6 UPDATE STATISTICS	43
4.8.7 USE database	43
4.9 InterSystems Extensions	43
4.9.1 OBJECTSCRIPT	44
4.9.2 IMPORTASQUERY	44
5 TSQL Settings	45
5.1 DIALECT	46
5.2 ANSI_NULLS	46

5.3 CASEINSCOMPARE	46
5.4 QUOTED_IDENTIFIER	47
5.5 Equal Literal Replacement	47
5.6 TRACE	48
6 TSQL Functions	51
6.1 Supported Functions	51
6.1.1 ABS	51
6.1.2 ACOS	51
6.1.3 ASCII	51
6.1.4 ASIN	51
6.1.5 ATAN	51
6.1.6 AVG	52
6.1.7 CAST	52
6.1.8 CEILING	52
6.1.9 CHAR	52
6.1.10 CHAR_LENGTH / CHARACTER_LENGTH	53
6.1.11 CHARINDEX	53
6.1.12 COALESCE	53
6.1.13 COL_NAME	53
6.1.14 CONVERT	53
6.1.15 COS	54
6.1.16 COT	54
6.1.17 COUNT	55
6.1.18 CURRENT_DATE	55
6.1.19 CURRENT_TIME	55
6.1.20 CURRENT_TIMESTAMP	55
6.1.21 CURRENT_USER	56
6.1.22 DATALENGTH	56
6.1.23 DATEADD	56
6.1.24 DATEDIFF	57
6.1.25 DATENAME	57
6.1.26 DATEPART	58
6.1.27 DAY	59
6.1.28 DB_NAME	59
6.1.29 DEGREES	59
6.1.30 ERROR_MESSAGE	59
6.1.31 ERROR_NUMBER	59
6.1.32 EXEC	59
6.1.33 EXP	59
6.1.34 FLOOR	60
6.1.35 GETDATE	60
6.1.36 GETUTCDATE	60
6.1.37 HOST_NAME	60
6.1.38 INDEX_COL	60
6.1.39 ISNULL	60
6.1.40 ISNUMERIC	60
6.1.41 LEFT	61
6.1.42 LEN	61
6.1.43 LOG	61
6.1.44 LOG10	61

6.1.45 LOWER	61
6.1.46 LTRIM	61
6.1.47 MAX	62
6.1.48 MIN	62
6.1.49 MONTH	62
6.1.50 NCHAR	62
6.1.51 NEWID	62
6.1.52 NOW	63
6.1.53 NULLIF	63
6.1.54 OBJECT_ID	63
6.1.55 OBJECT_NAME	63
6.1.56 PATINDEX	63
6.1.57 PI	64
6.1.58 POWER	64
6.1.59 QUOTENAME	64
6.1.60 RADIANS	64
6.1.61 RAND	65
6.1.62 REPLACE	65
6.1.63 REPLICATE	65
6.1.64 REVERSE	65
6.1.65 RIGHT	65
6.1.66 ROUND	65
6.1.67 RTRIM	66
6.1.68 SCOPE_IDENTITY	66
6.1.69 SIGN	66
6.1.70 SIN	66
6.1.71 SPACE	66
6.1.72 SQRT	67
6.1.73 SQUARE	67
6.1.74 STR	67
6.1.75 STUFF	67
6.1.76 SUBSTRING	67
6.1.77 SUM	68
6.1.78 SUSER_NAME	68
6.1.79 SUSER_SNAME	68
6.1.80 TAN	68
6.1.81 TEXTPTR	68
6.1.82 TEXTVALID	68
6.1.83 UNICODE	69
6.1.84 UPPER	69
6.1.85 USER	69
6.1.86 USER_NAME	69
6.1.87 YEAR	69
6.2 Unsupported Functions	69
7 TSQL Variables	71
7.1 Local Variables	71
7.1.1 Declaring a Local Variable	71
7.1.2 Setting a Local Variable	71
7.1.3 Initial and Default Values	72
7.1.4 Plain Local Variables	72

7.2 @@ Special Variables	72
7.2.1 @@ERROR	73
7.2.2 @@FETCH_STATUS	73
7.2.3 @@IDENTITY	73
7.2.4 @@LOCK_TIMEOUT	73
7.2.5 @@NESTLEVEL	74
7.2.6 @@ROWCOUNT	74
7.2.7 @@SERVERNAME	74
7.2.8 @@SPID	74
7.2.9 @@SQLSTATUS	75
7.2.10 @@TRANCOUNT	75
7.2.11 @@VERSION	75
8 TSQL System Stored Procedures	77
8.1 sp_addtype	77
8.2 sp_droptype	78
8.3 sp_procxmode (Sybase only)	78

About This Book

This book describes how to migrate schemas and stored procedures from Sybase or SQL Server and it will provide you with an understanding of the TSQL (Transact-SQL) implementation in InterSystems IRIS® data platform.

The book addresses a number of topics:

- An [Overview](#), which includes configuring TSQL, migrating source code and data, and dynamic execution of TSQL code.
- [InterSystems TSQL Constructs](#) including temporary tables, stored procedures, and transaction management.
- [InterSystems TSQL Language Elements](#): data types, operators, literals, reserved words.
- [InterSystems TSQL Commands](#)
- [InterSystems TSQL Settings](#)
- [InterSystems TSQL Functions](#)
- [InterSystems TSQL Variables](#)
- [InterSystems TSQL System Stored Procedures](#)

For a detailed outline, see the [Table of Contents](#).

When using InterSystems TSQL, you may find the following additional sources useful:

- *[The InterSystems SQL Reference](#)* provides details on individual SQL commands and functions, as well as information on the InterSystems SQL configuration settings, error codes, data types, and reserved words.
- *[Using the InterSystems SQL Gateway](#)* describes how to use the InterSystems SQL Gateway, which enables you to treat external tables as if they were native InterSystems IRIS® tables.
- *[Using the InterSystems ODBC Driver](#)* describes how to use InterSystems ODBC, which enables you to access InterSystems IRIS® tables via ODBC from external applications.
- *[Using Java with the InterSystems JDBC Driver](#)* describes how to use the InterSystems JDBC driver, which enables you to access InterSystems IRIS® tables via JDBC from external applications.

1

Overview

InterSystems TSQL is an implementation of Transact-SQL which supports many of the features of both the Microsoft and Sybase implementations. Transact-SQL is used with Microsoft SQL Server (MSSQL) and Sybase Adaptive Server.

InterSystems TSQL also contains a few proprietary extensions not found in either of these implementations. These are described in the [Commands](#) chapter.

Regardless of which InterSystems IRIS interface is used, TSQL code is used to generate corresponding IRIS SQL executable code. InterSystems IRIS does not provide system-level support for native TSQL.

This document will help you to quickly migrate schemas and stored procedures from Microsoft or Sybase databases and it will provide you with an understanding of the TSQL (Transact-SQL) implementation for InterSystems IRIS® data platform.

Microsoft provides good TSQL reference material at: <http://msdn.microsoft.com/en-us/library/bb545450.aspx>.

1.1 Migrating Existing TSQL Applications

To migrate existing TSQL applications to InterSystems TSQL, you need to perform three operations: configure InterSystems IRIS for TSQL, migrate the TSQL source code, and migrate the data.

1.1.1 Configuring TSQL

To configure your system for TSQL:

- Go into the InterSystems IRIS Management Portal. Select **System Administration, Configuration, SQL and Object Settings**, then select **TSQL Compatibility**. Here you can specify the dialect (Sybase or MSSQL), and turn on or off the **ANSI_NULLS**, **CASEINSCOMPARE**, and **QUOTED_IDENTIFIER** settings. The default for all three is “off”. These values are used to set the ^%SYS(“tsql”, “SET”, ...) global array values.
- From the Management Portal, select **System Administration, Configuration, SQL and Object Settings**, then **SQL**. From here, you can set the **Default Schema**. This is the [system-wide default schema name](#) (which maps to a package) for all unqualified DDL entities.
- From the Management Portal, select **System Administration, Configuration, SQL and Object Settings**, then **User DDL Mappings**. You can use this option to map any needed user-defined data types.

1.1.2 Migrating Source Code

The initial application migration is simple:

1. *Import the DDL:* Import table and view definitions using either the `%SYSTEM.SQL.DDLImport()` method (for single files) or the `%SYSTEM.SQL.DDLImportDir()` method (for multiple files in a directory). Within these methods, you set the `DDLMode` parameter to either "MSSQLServer" or "Sybase". These methods import DDL statements, as well as DML statements such as INSERT, convert them to equivalent InterSystems IRIS SQL, and execute them. For further details, see [Importing SQL Code](#).

Alternatively, you can invoke the `$SYSTEM.SQL.Sybase()` or `$SYSTEM.SQL.MSSQLServer()` method to import the schema. For further details, see [Importing SQL Code](#).

If the TSQL source contains CREATE PROC statements, then a class method containing the CREATE PROC source is created. InterSystems IRIS places this class method in either an existing class or in a new class whose name is based on the schema and procedure name. If the procedure already exists, then the existing version is replaced by the new version. If a class matching the class name generated from the schema and procedure already exists, then this class name is used — if it was previously generated by the TSQL utility. If not, then a unique class name is generated, based on the schema and procedure name. The resulting class is compiled once the procedure has been successfully created. If logging is requested then the source statements are logged along with the name of the containing class, class method, and the formal arguments generated. Any errors encountered by the process are also reported in the log. If an error is detected during CREATE PROC processing, InterSystems IRIS deletes any new class that was generated for that procedure.

2. *Inspect the log file for errors:* Search by Error #. A summary count of errors and successful imports will appear at the end of the log. In most cases, errors can be worked around or addressed by using information found in this document.
3. *Compile:* When you import DDL, table and view definition compilation is automatically performed. To compile other TSQL source code, it is best to use the command as follows:

```
DO $SYSTEM.OBJ.CompileAll("-l")
```

The lowercase "L" qualifier flag specifies that locking *is not* applied for the duration of the compile. For a full list of flag qualifiers, call `DO $SYSTEM.OBJ.ShowFlags()`.

1.1.3 Migrating the Data

In the Management Portal select **System Explorer**, **SQL**, then from the **Wizards** drop-down list select the **Data Migration Wizard**. This runs a wizard to migrate data from an external source and creates an InterSystems IRIS class definition to store it.

1.2 Writing and Executing TSQL on InterSystems IRIS

- *Writing TSQL class methods and stored procedures*

Create a class method stored procedure and enter the language as tsql. You can use the following template as a starting point:

```
ClassMethod MyTestMethod() As %Integer
[ Language = tsql, ReturnResultSets, SqlName=name, SqlProc ]
{
}
```

See the [Language](#) keyword for method definition in the *Class Definition Reference*.

You can write and maintain TSQL stored procedures (SPs) in Studio. A TSQL SP can be either a class method or a query. A class method takes parameters and returns a single scalar result, a query takes parameters and returns rows. If you put plain SELECT statements into a class method they will be executed but you won't be able to get the rows.

TSQL procedures are converted to InterSystems IRIS methods or queries with a Language type equal to TSQL. Use the following command:

```
DO ##class(%TSQL.Manager).load("sybase",<filename>,<logname>)
```

When compiling TSQL methods, ObjectScript code is generated. There is no system-level support for native TSQL. It is best to maintain the methods in TSQL to retain the familiar look of the original stored procedures.

- *Using Dynamic SQL*

InterSystems IRIS Dynamic SQL can be used to execute TSQL code queries and a limited subset of other DML and DDL statements.

- You can create a Dynamic SQL statement class instance, then set the [%Dialect property](#) to Sybase or MSSQL. You then prepare and execute a TSQL command within that object instance.
- You can execute Dynamic SQL without creating a statement class instance by invoking the [%SYSTEM.SQL.Prepare\(\)](#) method which prepares an SQL command, or the [%SYSTEM.SQL.Execute\(\)](#) method, which both prepares and executes an SQL command. Both of these methods provide a *Dialect* parameter.

See “[Using Dynamic SQL](#)” in the *Using InterSystems SQL* manual.

- *Using the Management Portal SQL Interface*

The SQL interface [Dialect option](#) allows you to set the SQL dialect to IRIS, Sybase, or MSSQL. The default is IRIS. Note that the dialect you select becomes the [user customized default](#) the next time you access the Management Portal. See “[Using the Management Portal SQL Interface](#)” in the *Using InterSystems SQL* manual.

- *Using the TSQL Shell*

The InterSystems TSQL Shell can be used to execute Transact-SQL code from InterSystems IRIS. To use the TSQL Shell, invoke the [TSQLShell\(\)](#) (or [TSQL\(\)](#)) method from the Terminal as follows: `DO $SYSTEM.SQL.TSQLShell()`. This invokes the [InterSystems SQL Shell](#) and sets its DIALECT configuration parameter to the [currently configured TSQL dialect](#) (MSSQL or Sybase). The initial configuration default is MSSQL.

When entering SQL code interactively, the TSQL Shell supports, but does not require, the semicolon (;) statement delimiter at the end of each SQL statement.

You can use the Shell’s [RUN command](#) to execute a TSQL script file. The RUN command displays a series of prompts, including `Please enter the end-of-statement delimiter (Default is 'GO'):` `GO=>`. This enables you to specify the TSQL semicolon (;) as the statement delimiter in your script file, rather than the InterSystems IRIS default GO statement. See “[Using the SQL Shell Interface](#)” in the *Using InterSystems SQL* manual.

- *Using the InterSystems SQL Shell*

The InterSystems SQL Shell can be used to execute lines of TSQL code by using the [SET DIALECT](#) command to set the Shell’s dialect to Sybase or MSSQL.

When the Shell’s dialect is set to Sybase or MSSQL, the SQL Shell supports, but does not require, the semicolon (;) statement delimiter at the end of each SQL statement. When the Shell’s dialect is set to IRIS, a semicolon (;) statement delimiter results in an SQLCODE -25 error.

You can use the Shell’s [RUN command](#) to execute a TSQL script file. The RUN command displays a series of prompts, including `Please enter the end-of-statement delimiter (Default is 'GO'):` `GO=>`. This enables you to specify the TSQL semicolon (;) as the statement delimiter in your script file, rather than the InterSystems IRIS default GO statement. See “[Using the SQL Shell Interface](#)” in the *Using InterSystems SQL* manual.

- *Using Triggers*

You can write and maintain triggers, which are sets of instructions that appear in TSQL code and that are executed in response to certain SQL events. You can use the [Language=tsql](#) class definition keyword to specify that a trigger is

written in TSQL. The [UpdateColumnList](#) class definition keyword is only supported for TSQL. Row-level triggers are not supported for TSQL. See “[Using Triggers](#)” in the *Using InterSystems SQL* manual.

2

InterSystems TSQL Constructs

2.1 Table References

InterSystems TSQL supports table references with the InterSystems IRIS® data platform SQL format:

```
schema.table
```

The only mandatory table reference component is `table`. If the schema is omitted, TSQL uses the [system-wide default schema name](#).

Other forms of Transact-SQL may use table references with up to four components, separated by dots. Here is how a Transact-SQL table reference is processed:

- The `server.` prefix, if present, is ignored.
- The `database.` prefix, if present, is removed. InterSystems IRIS only supports one database name: 'master'.
- The `user.` prefix, if present, is mapped to the *schema* name.

For the purposes of name translation, a field name has the field suffix removed while translation is performed and then replaced afterwards.

2.2 Temporary Tables

InterSystems TSQL supports `#tablename` temporary tables. A `#tablename` temporary table is visible to the current procedure of the current process. It is also visible to any procedure called from the current procedure. `#tablename` syntax is only supported in TSQL procedures (class methods projected as procedures with language `tsql`).

A temporary table is defined by using **CREATE TABLE** with a table name starting with "#". The temporary table is created at runtime. A `#tablename` table definition goes out of scope when you exit the procedure. All temporary table definitions go out of scope when the connection is dropped. You can also explicitly delete a temporary table using **DROP TABLE**.

However, if a temporary table is referenced by an active result set, the temporary table may become invisible to the process, but the data and definition are retained until the result set goes out of scope.

A `#tablename` temporary table is visible both to the creating procedure and to any procedures called from that procedure. Temporary tables are visible to nested procedure calls. It is not necessary to declare the temporary table in the called procedure. If the called procedure also creates a temporary table with the same name, InterSystems IRIS uses the most recently

created table definition. Because a temporary table is defined using an ObjectScript local variable, the creation, modification, and deletion of these tables are not journaled transaction events; rolling back the transaction has no effect on these operations.

2.3 System Tables

System tables exist per InterSystems IRIS namespace.

Systypes

Partially supported.

2.4 Transactions

Code generated for **BEGIN TRAN**, **COMMIT** and **ROLLBACK** uses explicit transaction mode, but following a transaction TSQL always restores the mode which was active before the **BEGIN TRAN** statement. TSQL restores this mode when the procedure is exited from, or when a **COMMIT** or **ROLLBACK** is issued, whichever comes first.

2.5 Cursor Name Management

You can declare the same cursor more than once, so long as only one version of the cursor is open at runtime. If the same cursor is declared more than once in a stored procedure, all but the first declaration are associated with renamed cursors. **OPEN**, **FETCH**, **CLOSE**, and **DEALLOCATE** statements are assumed to refer to the most recent **DECLARE** for the given cursor. Note that the lexical position of a statement within a stored procedure is all that is used to match up a cursor name with its **DECLARE** — no account is taken of runtime paths through the code.

Cursors inside queries are named using an extension of the scheme used in InterSystems SQL queries. For example:

```
DECLARE C CURSOR FOR SELECT A FROM B
--
OPEN C
FETCH C
CLOSE C
DEALLOCATE C
--
DECLARE C CURSOR FOR SELECT D FROM E
--
OPEN C
FETCH C
CLOSE C
DEALLOCATE C
```

Would be effectively translated to:

```
DECLARE C CURSOR FOR SELECT A FROM B
--
OPEN C
FETCH C
CLOSE C
DEALLOCATE C
--
DECLARE Cv2 CURSOR FOR SELECT D FROM E
--
OPEN Cv2
FETCH Cv2
CLOSE Cv2
DEALLOCATE Cv2
```


2.6 SYSOBJECTS References

Commonly, an application will have setup procedures that create tables, views, and the metadata for the application environment. Such procedures will have expressions like:

```
IF EXISTS (SELECT * FROM SYSOBJECTS
WHERE ID = OBJECT_ID('People'))
```

This determines if a table exists, in this example. It's usually followed by a DROP and CREATE statement to reestablish the table metadata.

TSQL procedures and triggers can reference the SYSOBJECTS system table. InterSystems TSQL supports the following columns in the SYSOBJECTS table (%TSQL.sys.objects class properties):

<i>name</i>	Object name.
<i>id</i>	Object Id.
<i>type</i>	Object type: can be one of the following values: K=PRIMARY KEY or UNIQUE constraint; P=stored procedure; RI=FOREIGN KEY constraint; S=system table; TR=trigger; U=user table; V=view.
<i>deltrig</i>	Object ID of a delete trigger if the entry is a table. Table ID of a table if the entry is a trigger.
<i>instrig</i>	Object ID of a table's insert trigger if the entry is a table.
<i>updtrig</i>	Object ID of a table's update trigger if the entry is a table.
<i>parent_obj</i>	Object identification number of parent object. For example, the table ID if a trigger or constraint.
<i>schema</i>	Name of the schema in which the object resides.
<i>parent_obj_name</i>	Object name of parent_obj. If <i>parent_obj</i> =0, <i>parent_obj_name</i> is the same as <i>name</i> .

The SYSOBJECTS table is read-only. The SYSOBJECTS table may be referenced from outside a TSQL procedure or trigger by the name %TSQL_sys.objects. SYSOBJECTS is not supported for tables mapped across namespaces.

Note: InterSystems IRIS provides the %Dictionary package of class objects that can perform the same operations as SYSOBJECTS references. For further details, refer to the %Dictionary package in the *InterSystems Class Reference*.

3

InterSystems TSQL Language Elements

This chapter describes the following TSQL language elements for InterSystems IRIS® data platform:

- [Literals, Reserved Words, and Comments](#)
- [Identifiers](#)
- [Data Types](#)
- [Arithmetic, Comparison, String, Logical, and Bitwise Operators](#)

3.1 Literals

3.1.1 String Literals

A string literal must be delimited by quote characters. The preferred delimiter characters are single quote characters. You can also use double quote characters as string delimiters if you specify `SET DELIMITED_IDENTIFIER OFF`. Otherwise, double quote characters are parsed as delimiting an identifier.

If you delimit a string literal with single quote characters, you can include literal double quote characters within the string. To include a literal single quote character within the string, double it by typing two single quotes.

A string containing literal single quotes, such as `'this is an ''embedded'' string'`, is compiled by InterSystems IRIS to single quotes within double quotes: `"this is an 'embedded' string"`.

3.1.2 Empty Strings

When migrating Transact-SQL code to InterSystems TSQL, it may be necessary to redefine the empty string. You can do this by setting the following InterSystems IRIS system global:

```
^%SYS("sql","sys","namespace",nspace,"empty string")
```

All of these specified values are keyword literals, except *nspace*, which is a namespace name specified as a quoted string.

CAUTION: Changing the empty string definition should be done with extreme caution. It can result in data containing different representations for an empty string. It can also cause existing programs to fail when executed in this namespace. After defining the empty string, you must purge all cached queries and recompile all classes and routines for that namespace that use the former empty string definition.

The following ObjectScript example changes the empty string definition for the SAMPLES namespace. It first sets the empty string value to a single blank space. It then sets the empty string value to the non-printing character represented by the ASCII code 0. (This example then immediately resets the empty string value to the InterSystems IRIS default):

```
SET ^%SYS("sql","sys","namespace","SAMPLES","empty string")=" "
WRITE !,"Empty string set to:"
ZZDUMP ^%SYS("sql","sys","namespace","SAMPLES","empty string")
SET ^%SYS("sql","sys","namespace","SAMPLES","empty string")=$CHAR(0)
WRITE !,"Empty string set to:"
ZZDUMP ^%SYS("sql","sys","namespace","SAMPLES","empty string")
SET ^%SYS("sql","sys","namespace","SAMPLES","empty string")=""
WRITE !,"Empty string reset to:"
ZZDUMP ^%SYS("sql","sys","namespace","SAMPLES","empty string")
WRITE !,!, "End of sample program"
```

3.1.3 NULL

In TSQL a NULL supplied to a boolean operation returns as FALSE, as shown in the following example:

```
DECLARE @var BINARY(1)
SELECT @var=NULL
IF @var PRINT "true" ELSE PRINT "false"
```

In Sybase dialect, NULL is equal to NULL. A NULL=NULL comparison returns TRUE, and a NULL != NULL comparison returns FALSE.

In MSSQL dialect, a comparison of NULL with any value returns FALSE. Thus NULL=NULL and NULL != NULL comparisons both return FALSE.

```
DECLARE @var BINARY(1)
SELECT @var=NULL
IF @var=NULL PRINT "true" ELSE PRINT "false"
```

In Sybase dialect, NULL is not equal to any value. Therefore, Not Equals (!=) comparison involving NULL and any boolean, numeric, or string value (including the empty string ("")) returns TRUE. All Equals (=), Greater Than (>) or Less Than (<) comparisons return FALSE.

In MSSQL dialect, NULL cannot be compared to a value. Thus all Equals (=), Not Equals (!=), Greater Than (>) or Less Than (<) comparisons return FALSE.

In a TSQL string concatenation operation, NULL is equivalent to an empty string. In a TSQL arithmetic operation, NULL is equivalent to 0.

3.1.4 Hexadecimal

InterSystems TSQL automatically converts hexadecimal numeric literals in TSQL source code to the corresponding decimal (base-10) numeric literals.

3.1.5 Reserved Words

InterSystems TSQL cannot use as identifiers the SQL Server reserved words. InterSystems TSQL can use InterSystems SQL reserved words (that are not also SQL Server reserved words) if the **QUOTED_IDENTIFIER** SQL configuration setting is set to Yes.

3.1.6 Comments, Blank Lines, and Semicolons

InterSystems TSQL supports both single-line and multi-line comments.

- A single line comment continues to the rest of the line. When used in the TSQL shell, a comment does not encompass the end-of-line qualifier, such as /x or /c. InterSystems TSQL supports both — and // as single-line comment delimiters.

- A multi-line comment begins with `/*` and ends with `*/`. A comment can include nested `/* ... */` comments.

```
PRINT 'these are comments'
-- this is a single-line comment
// this is a single-line comment
/* This is a multi-line comment
The command
PRINT 'do not print'
is part of the comment and is not executed */
```

3.1.6.1 TSQL-only Statements

InterSystems TSQL provides the means to include executable statements within InterSystems IRIS TSQL code which are parsed as nonexecutable comments in Transact-SQL. A statement prefixed with two hyphens and a vertical bar is parsed by InterSystems IRIS as an executable statement. Sybase Adaptive Server and Microsoft SQL Server consider this to be a Transact-SQL comment.

```
PRINT 'any context'
-- PRINT 'commented out'
--| PRINT 'InterSystems only'
```

3.1.6.2 Semicolons

You can specify a blank line by using either two hyphens or a semicolon.

A semicolon either before or after a TSQL statement is ignored. They are supported for compatibility with Transact-SQL code, such as stored procedures, that ends statements with a semicolon.

```
PRINT 'no semicolon'
--
PRINT 'trailing semicolon';
;
;PRINT 'leading semicolon'
```

3.2 Identifiers

An identifier is a name for a TSQL object, such as a table, column, view, key, index, trigger, or stored procedure. Naming conventions for identifiers are as follows:

- The first character of an identifier must be a letter, an underscore (`_`) or a percent (`%`) character.
- Subsequent characters of an identifier may be letters, numbers, underscores (`_`), dollar signs (`$`), or pound signs (`#`).
- Identifiers can be of any length, but must be unique within their first 30 characters.
- Identifiers are not case-sensitive.
- An identifier cannot be an SQL reserved word.
- A pound sign (`#`) prefix to an identifier specifies that it is the name of a temporary table.
- An at sign (`@`) prefix to an identifier specifies that it is the name of a variable.

Some identifiers are qualified with a schema name. For example, `schema.tablename` or `schema.storedprocedure`. If the schema name is omitted, the identifier is unqualified. TSQL resolves unqualified identifiers by using either the [system-wide default schema](#) (for DDL) or the `schemaPath` property (for DML), which provides a search path of schemas to check for the specified table name or stored procedure name.

3.2.1 Delimited and Quoted Identifiers

A delimited identifier is not restricted by the naming conventions of ordinary identifiers. For example, a delimited identifier can be the same word as an SQL reserved word; a delimited identifier can contain space characters.

By default, both square brackets and double quotation marks can be used to delimit an identifier. These delimiters are interchangeable; you can define a delimited identifier by enclosing it with square brackets, and invoke the same delimited identifier by specifying it enclosed with double quotation marks.

You can specify a quoted identifier if the **QUOTED_IDENTIFIER** SQL configuration setting is set to Yes. You specify a quoted identifier by enclosing it in double quotation marks. When **QUOTED_IDENTIFIER** is on, double quotes are parsed as delimiting an identifier. When **QUOTED_IDENTIFIER** is off, double quotes are parsed as alternative delimiters for string literals. The preferable delimiters for string literals are single quotes. A quoted identifier can contain any characters, including blank spaces.

3.3 Data Types

The following data types are supported for local variables and table columns. These data types are supported in that they are parsed as valid data types; however, no range or value validation is performed.

BINARY(*n*) and **VARBINARY**(*n*). The (*n*) size specification is mandatory.

BIT

BOOLEAN

CHAR and **VARCHAR**

CHAR(*n*), **NCHAR**(*n*), **VARCHAR**(*n*), and **NVARCHAR**(*n*)

VARCHAR(MAX), and **NVARCHAR**(MAX). By default, these map to %Stream.GlobalCharacter.

DATETIME and **SMALLDATETIME**

DECIMAL, **DECIMAL**(*p*), and **DECIMAL**(*p*,*s*). Where *p* and *s* are integers specifying precision (total digits) and scale (decimal digits).

DOUBLE and **DOUBLE PRECISION**

FLOAT and **FLOAT**(*n*)

INT, **BIGINT**, **SMALLINT**, and **TINYINT**

MONEY and **SMALLMONEY**

NATIONAL

NUMERIC, **NUMERIC**(*p*), and **NUMERIC**(*p*,*s*). Where *p* and *s* are integers specifying precision (total digits) and scale (decimal digits).

REAL

TIMESTAMP

Note: The Microsoft SQL Server **TIMESTAMP** data type is *not* used for date or time information. It is an integer counter of the number of times a record is inserted or updated in a table. It should not be confused with the InterSystems SQL and ODBC **TIMESTAMP** data type, which represents a date and time in YYYY-MM-DD HH:MM:SS.nnnnnnnn format. In TSQL, use **DATETIME** and **SMALLDATETIME** for date and time values.

ROWVERSION

SQL_VARIANT

The following SQL Server data types are supported in a specific context:

CURSOR

NTEXT, TEXT By default, these map to %Stream.GlobalCharacter.

IMAGE

TABLE

The following are not implemented:

- UNIQUEIDENTIFIER stored as a 16-byte binary string. Instead use VARCHAR(32) as the data type for a globally unique ID.
- SQL92 and TSQL options
- UPDATE OF

3.4 Operators

3.4.1 Arithmetic and Equality Operators

InterSystems TSQL supports + (addition), – (subtraction), * multiplication, / division, and % modulo arithmetic operators.

InterSystems TSQL supports the following equality and comparison operators:

- = (equal to)
- <> (not equal to) and != (not equal to)
- < (less than), !< (not less than), <= (less than or equal to)
- > (greater than), !> (not greater than), >= (greater than or equal to)

When performing equality comparisons (= or <>) between date values with different data types, all date and time values are compared using the TIMESTAMP data type. Thus two dates in different formats can be meaningfully compared. A date value declared as a STRING data type can be compared to a date value declared as a DATETIME data type.

3.4.2 Concatenation Operator

InterSystems TSQL supports the + (plus sign) as both a concatenation operator and the addition operator. The plus sign functions as a concatenation operator with strings. You can concatenate several strings together using this operator. If all items are strings, TSQL performs concatenation; however, if one of the items is a number, TSQL performs addition, treating non-numeric strings as 0.

'world'+'wide'+'web' concatenates to 'worldwideweb'

'world'+'33'+'web' concatenates to 'world33web'

'world'+33+'web' performs addition (0+33+0=33)

In a TSQL string concatenation operation, NULL is equivalent to an empty string. In a TSQL arithmetic operation, NULL is equivalent to 0. Note that because the plus sign (+) is used for both concatenation and addition, the data type declaration of the NULL variable is critical. The following examples all return “bigdeal”:

```
DECLARE @var1 BINARY(1)
DECLARE @var2 VARCHAR(10)
SELECT @var1=NULL,@var2=NULL
PRINT "big"+NULL+"deal"
PRINT "big"+@var1+"deal"
PRINT "big"+@var2+"deal"
```

The following example returns 0; it treats the + as an arithmetic operator and interprets the argument as $0 + 0 + 0 = 0$:

```
DECLARE @var1 INT
SELECT @var1=NULL
PRINT "big"+@var1+"deal"
```

InterSystems TSQL also supports || as a concatenation operator.

3.4.3 Comparison Operators

3.4.3.1 BETWEEN

InterSystems TSQL supports the **BETWEEN** range check operator of the form: `BETWEEN num1 AND num2`. **BETWEEN** is inclusive of the specified range limits.

3.4.3.2 IS NULL

InterSystems TSQL supports the IS NULL match operator. A variable is NULL if it has been declared but not assigned a value, or if it has been explicitly specified as NULL. The empty string is not NULL.

3.4.3.3 LIKE

InterSystems TSQL supports the LIKE pattern match operator. LIKE performs not case-sensitive matching of letters. InterSystems TSQL also supports NOT LIKE.

3.4.4 NOT Logical Operator

The NOT logical operator inverts the truth value of the statement that follows it. For example, `IF NOT EXISTS(. . .)`. NOT is not case-sensitive.

3.4.5 Bitwise Logical Operators

InterSystems TSQL supports the AND (&), OR (|), XOR (^), and NOT (~) bitwise operators for the integer data type. The decimal integers are converted to binary, the logical operation is performed, and the resulting binary is converted to a decimal integer value. The NOT (~) operator is a unary operator that inverts bits.

4

TSQL Commands

This chapter lists the supported TSQL commands for InterSystems IRIS® data platform in the following groups:

- Data Definition Language (DDL) statements:
ALTER TABLE, CREATE TABLE, DROP TABLE
CREATE INDEX, DROP INDEX
CREATE TRIGGER, DROP TRIGGER
CREATE VIEW, DROP VIEW
Parsed but ignored: CREATE DATABASE, DROP DATABASE
- Data Management Language (DML) statements:
INSERT, UPDATE, DELETE, TRUNCATE TABLE
READTEXT, WRITETEXT, UPDATETEXT
- Query statements:
SELECT, JOIN, UNION, FETCH cursor
- Flow of control statements:
IF, WHILE, CASE, GOTO, WAITFOR
- Assignment statements:
DECLARE, SET
- Transaction statements:
SET TRANSACTION ISOLATION LEVEL, BEGIN TRANSACTION, COMMIT, ROLLBACK, LOCK TABLE
Parsed but ignored: SAVE TRANSACTION
- Procedure statements
CREATE PROCEDURE, DROP PROCEDURE
CREATE FUNCTION, ALTER FUNCTION, DROP FUNCTION
RETURN, EXECUTE, CALL
- Other statements
CREATE USER, GRANT, REVOKE, PRINT, RAISERROR, UPDATE STATISTICS
- InterSystems IRIS extensions

OBJECTSCRIPT, IMPORTASQUERY

InterSystems IRIS implementation of TSQL accepts, but does not require, a semicolon command terminator. When importing TSQL code to InterSystems SQL, semicolon command terminators are stripped out.

4.1 Data Definition Language (DDL) Statements

The following DDL statements are supported.

4.1.1 CREATE TABLE

Defines a table, its fields, and their data types and constraints.

```
CREATE TABLE [schema. | #]tablename (fieldname datatype constraint [...])
```

CREATE TABLE can specify a table using table (default schema), *schema.table*, or *database..table* syntax. See [Table References](#).

A **CREATE TABLE** can create a [temporary table](#) by prefixing a # character to the table name. A temporary table can only be defined from a stored procedure; you cannot define a temporary table from Dynamic SQL outside of a stored procedure. To create a fully-qualified temporary table name, use quotes around each name element such as the following:

```
"SQLUser"."#mytemp".
```

A valid table name must begin with a letter, an underscore character (`_`), or a # character (for a local temporary table). Subsequent characters of a table name may be letters, numbers, or the #, \$, or `_` characters. Table names are not case-sensitive.

A field name must be a valid [TSQL identifier](#). A field name can be delimited using square brackets. This is especially useful when defining a field that has the same name as a reserved word. The following example defines two fields named Check and Result:

```
CREATE TABLE mytest ([Check] VARCHAR(50),[Result] VARCHAR(5))
```

The optional **CONSTRAINT** keyword can be used to specify a user-defined constraint name for a field constraint or a table constraint. You can specify multiple **CONSTRAINT name type** statements for a field.

InterSystems SQL does not retain constraint names. Therefore these names cannot be used by a subsequent **ALTER TABLE** statement.

The table field constraints **DEFAULT**, **IDENTITY**, **NULL**, **NOT NULL**, **PRIMARY KEY**, **[FOREIGN KEY] REFERENCES** (the keywords **FOREIGN KEY** are optional), **UNIQUE**, **CLUSTERED**, and **NONCLUSTERED** are supported. The table constraint **FOREIGN KEY REFERENCES** is supported.

The field definition **DEFAULT** values can include the following TSQL functions: **CURRENT_TIMESTAMP**, **CURRENT_USER**, **GETDATE**, **HOST_NAME**, **ISNULL**, **NULLIF**, and **USER**.

The field definition **IDENTITY** constraint is supported and assigned a system-generated sequential integer. The **IDENTITY** arguments *seed* and *increment* are parsed, but ignored.

The TSQL **CREATE TABLE** command can create a [sharded table](#). The syntax for the **SHARD** clause is the same as for the [InterSystems SQL CREATE TABLE](#) statement:

```
SHARD [ KEY fieldname { , fieldname2 } ] [ COSHARD [ WITH ] [( tablename )] ]
```

The CHECK field constraint is not supported. If a CHECK constraint is encountered while compiling TSQL source Inter-Systems IRIS generates an error message indicating that CHECK constraints are not supported. This error is logged in the compile log (if active), and the source is placed in the unsupported log (if active).

If the table already exists, an SQLCODE -201 error is issued.

The following Dynamic SQL example creates a temporary table named #mytest with four fields, populates it with data, then displays the results. The LastName field has multiple constraints. The FirstName field takes a default. The DateStamp field takes a system-defined default:

```
SET sql=9
SET sql(1)="CREATE TABLE #mytest (MyId INT PRIMARY KEY,"
SET sql(2)="LastName VARCHAR(20) CONSTRAINT unq_lname UNIQUE "
SET sql(3)="CONSTRAINT nonull_lname NOT NULL,"
SET sql(4)="FirstName VARCHAR(20) DEFAULT '***TBD***',"
SET sql(5)="DateStamp DATETIME DEFAULT CURRENT_TIMESTAMP)"
SET sql(6)="INSERT INTO #mytest(MyId,LastName,FirstName) VALUES (1224,'Smith','John') "
SET sql(7)="INSERT INTO #mytest(MyId,LastName) VALUES (1225,'Jones') "
SET sql(8)="SELECT MyId,FirstName,LastName,DateStamp FROM #mytest"
SET sql(9)="DROP TABLE #mytest"
SET statement=##class(%SQL.Statement).%New()
SET statement.%Dialect="MSSQL"
SET status=statement.%Prepare(.sql)
WRITE status,!
SET result=statement.%Execute()
DO result.%Display()
```

4.1.1.1 Parsed But Ignored

The table constraint clauses WITH, ON, and TEXTIMAGE ON are parsed for compatibility, but are ignored. The *index_options* clause for the UNIQUE or PRIMARY KEY constraint is parsed for compatibility, but is ignored.

The following SQL Server parenthesized WITH options in a table constraint are parsed but ignored:

ALLOW_PAGE_LOCKS, ALLOW_ROW_LOCKS, DATA_COMPRESSION, FILLFACTOR, IGNORE_DUP_KEY, PAD_INDEX, and STATISTICS_NORECOMPUTE.

The field constraints CLUSTERED and NONCLUSTERED are parsed for compatibility, but are ignored.

4.1.2 ALTER TABLE

Modifies the definition of a table, its fields, and their data types and constraints.

The following syntactical forms are supported:

```
ALTER TABLE tablename ADD fieldname datatype [DEFAULT value]
    [{UNIQUE | NOT NULL} | CONSTRAINT constraintname {UNIQUE | NOT NULL} ]
ALTER TABLE tablename ALTER COLUMN fieldname newdatatype
ALTER TABLE tablename DROP COLUMN fieldname [,fieldname2]
ALTER TABLE tablename ADD tableconstraint FOR fieldname
ALTER TABLE tablename DROP tableconstraint
ALTER TABLE tablename DROP FOREIGN KEY role
ALTER TABLE tablename ADD CONSTRAINT constraint DEFAULT defaultvalue FOR fieldname
ALTER TABLE tablename ADD CONSTRAINT constraint FOREIGN KEY
ALTER TABLE tablename DROP CONSTRAINT constraint
```

Specify *tablename* as described in [Table References](#).

- **ALTER TABLE...ADD *fieldname*** can add a field definition or a comma-separated list of field definitions:
 - DEFAULT is supported.
 - NOT NULL is supported if the table contains no data. If the table contains data, you can only specify NOT NULL if the field also specifies a DEFAULT value.
 - UNIQUE is parsed but ignored. To establish a unique constraint use the CREATE INDEX command with the UNIQUE keyword.

The full supported syntax for **ALTER TABLE...ADD *fieldname*** is as follows:

```
ALTER TABLE tablename
  [ WITH CHECK | WITH NOCHECK ]
  ADD fieldname datatype [DEFAULT value]
  [{UNIQUE | NOT NULL} | CONSTRAINT constraintname {UNIQUE | NOT NULL} ]
  [ FOREIGN KEY (field1[,field2[,...]])
    REFERENCES tablename(field1[,field2[,...]]) ]
```

WITH CHECK | WITH NOCHECK is parsed by InterSystems IRIS, but is ignored. In Transact-SQL, WITH CHECK | WITH NOCHECK provides an execution time check of existing data for a new or newly enabled constraint. InterSystems TSQL does not specifically support that, although InterSystems SQL will check existing data against a new constraint.

The Sybase PARTITION BY clause is not supported.

- **ALTER TABLE...ALTER COLUMN *fieldname datatype*** can change the data type of an existing field. This command completes without error when the specified *datatype* is the same as the field's existing data type.
- **ALTER TABLE...DROP [COLUMN] *fieldname*** can drop a defined field or a comma-separated list of defined fields. The keyword **DELETE** is a synonym for the keyword **DROP**.
 - Sybase: the COLUMN keyword is not permitted, the CONSTRAINT keyword is required: **ALTER TABLE...DROP *fieldname*, CONSTRAINT *constraint***
 - MSSQL: the COLUMN keyword is required, the CONSTRAINT keyword is optional: **ALTER TABLE...DROP COLUMN *fieldname*, *constraint***
- **ALTER TABLE...DROP [CONSTRAINT] *constraintname*** can drop a constraint from a field. The keyword **DELETE** is a synonym for the keyword **DROP**.
 - Sybase: the CONSTRAINT keyword is required.
 - MSSQL: the CONSTRAINT keyword is optional.
- **ALTER TABLE...ADD CONSTRAINT...DEFAULT** syntax does not create a field constraint. Instead, it performs the equivalent of an **ALTER TABLE...ALTER COLUMN...DEFAULT** statement. This means that InterSystems IRIS establishes the specified field default as the field property's initial expression. Because no field constraint is defined, this "constraint" cannot be subsequently dropped or changed.

CHECK | NOCHECK CONSTRAINT is not supported by InterSystems IRIS TSQL. Specifying this CHECK or NOCHECK keyword generates an error message.

4.1.3 DROP TABLE

Deletes a table definition.

```
DROP TABLE [IF EXISTS] tablename
```

Deletes a table definition. You can delete both regular tables and temporary tables. (Temporary table names begin with a '#' character.) **DROP TABLE** ignores a nonexistent temporary table name and completes without error.

Specify *tablename* as described in [Table References](#).

If *tablename* has an associated view, you must delete the view before you can delete the table.

The IF EXISTS clause is parsed but ignored.

4.1.4 CREATE INDEX

Creates an index for a specified table or view.

```
CREATE [UNIQUE] INDEX indexname ON tablename (fieldname [,fieldname2])
```

You can create an index on a field or a comma-separated list of fields.

You can create an index on the IDKEY (which is treated as a clustered index), on an IDENTITY field (which create an index on the %%ID field), on the Primary Key, or on other fields.

Specify *tablename* as described in [Table References](#).

The UNIQUE keyword creates a unique value constraint index for the specified field(s).

The following Transact-SQL features are parsed, but ignored:

- The CLUSTERED/NONCLUSTERED keywords. Other than the IDKEY, which is implicitly treated as a clustered index, InterSystems TSQL does not support clustered indices.
- The ON *dbspace* clause.
- The ASC/DESC keywords.
- The INCLUDE clause.
- WITH clause options, such as WITH FILLFACTOR=*n* or WITH DROP_EXISTING=ON. The comma-separated list of WITH clause options can optionally be enclosed in parentheses.
- The ON filegroup or IN *dbspace-name* clause.

The following Transact-SQL features are not currently supported:

- Sybase index types.
- The IN *dbspace* clause.
- The NOTIFY *integer* clause.
- The LIMIT *integer* clause.
- Using a function name as an alternative to a field name.

The ALTER INDEX statement is not supported.

4.1.5 DROP INDEX

Deletes an index definition. You can delete a single index or a comma-separated list of indices, using either of the following syntax forms:

```
DROP INDEX tablename.indexname [,tablename.indexname]  
  
DROP INDEX indexname ON tablename [WITH (...)] [,indexname ON tablename [WITH (...)]  
]
```

tablename is the name of the table containing the indexed field. Specify *tablename* as described in [Table References](#).

indexname is the name of the index. It can be a [regular identifier](#) or a [quoted identifier](#).

The WITH (...) clause, with any value within the parentheses, is accepted by syntax checking for compatibility, but is not validated and performs no operation.

The IF EXISTS clause is not supported.

4.1.6 CREATE TRIGGER

Creates a statement-level trigger.

```
CREATE TRIGGER triggername ON tablename
[WITH ENCRYPTION]
{FOR | AFTER | INSTEAD OF} {INSERT | DELETE | UPDATE}
[WITH APPEND]
[NOT FOR REPLICATION]
AS tsql_trigger_code
```

You can create a trigger for one event (INSERT), or for a comma-separated list of events (INSERT,UPDATE).

Specify *tablename* as described in [Table References](#).

The FOR, AFTER, and INSTEAD OF keywords are synonyms. A trigger is always pulled after the event operation is performed.

If there are multiple triggers for the same event or comma-separated list of events they are executed in the order the triggers were created.

The following clauses are parsed but ignored: WITH ENCRYPTION, WITH APPEND, NOT FOR REPLICATION.

InterSystems TSQL does not support row-level triggers.

You cannot include a CREATE TRIGGER statement in [CREATE PROCEDURE](#) code.

4.1.7 DROP TRIGGER

Deletes a trigger definition.

```
DROP TRIGGER [owner.]triggername
```

4.1.8 CREATE VIEW

Creates a view definition.

```
CREATE VIEW [owner.]viewname
[WITH {ENCRYPTION | SCHEMABINDING | VIEW_METADATA}]
AS select_statement
[WITH CHECK OPTION]
```

A *viewname* must be a unique [TSQL identifier](#). Specify *viewname* as described in [Table References](#). If the view already exists, an SQLCODE -201 error is issued. A *viewname* can be a [delimited identifier](#). For example, CREATE VIEW Sample.[Name/Age View].

By default, the view fields have the same names as the fields in the SELECT table. To specify different names for the view fields, specify field aliases in the SELECT statement. These aliases are used as the view field names:

```
CREATE VIEW NameAgeV
AS SELECT Name AS FullName, Age AS Years FROM Sample.Person
```

You can specify a WITH clause with a single keyword or a comma-separated list of keywords. For example: WITH SCHEMABINDING, ENCRYPTION, VIEW_METADATA. The ENCRYPTION, SCHEMABINDING, and VIEW_METADATA keywords are parsed but ignored.

The *select_statement* can only include an ORDER BY clause if this clause is paired with a TOP clause. If you wish to include all of the rows in the view, you can pair an ORDER BY clause with a TOP ALL clause. You can include a TOP

clause without an ORDER BY clause. However, if you include an ORDER BY clause without a TOP clause, an SQLCODE -143 error is generated.

The *select_statement* can contain a UNION or UNION ALL.

The optional WITH CHECK OPTION clause prevents an update through the view that makes the record inaccessible to that view. It does this by checking the WITH clause in the SELECT statement. WITH CHECK OPTION binds to InterSystems SQL using the default of CASCADE.

The ALTER VIEW statement is not supported.

4.1.9 DROP VIEW

Deletes a view definition.

```
DROP VIEW viewname [,viewname2 [...] ]
```

You can delete a single view, or a comma-separated list of views. Specify *viewname* as described in [Table References](#).

DROP VIEW is not an all-or-nothing operation. It deletes existing views in the list of views until it encounters a nonexistent view in the list. At that point the delete operation stops with an SQLCODE -30 error.

The IF EXISTS clause is not supported.

4.1.10 CREATE DATABASE

CREATE DATABASE syntax is parsed to provide compatibility. No functionality is provided.

```
CREATE DATABASE dbname
```

Only this basic CREATE DATABASE syntax is parsed.

Sybase additional CREATE DATABASE clauses are not supported.

MSSQL attach a database and create a database snapshot syntax options are not supported.

The ALTER DATABASE statement is not supported.

4.1.11 DROP DATABASE

DROP DATABASE syntax is parsed to provide compatibility. No functionality is provided.

```
DROP DATABASE dbname
```

4.2 Data Management Language (DML) Statements

- TSQL can resolve an unqualified table name using a [schema search path](#) for a single DML statement in Dynamic SQL.
- TSQL *cannot* resolve an unqualified table name using a schema search path for multiple DML statements in Dynamic SQL. This includes multiple statements such as an explicit BEGIN TRANSACTION followed by a single DML statement.

4.2.1 DELETE

Deletes rows of data from a table. Both **DELETE** and **DELETE ... FROM** are supported:

```
DELETE FROM tablename WHERE condition

DELETE FROM tablename FROM matchtablename WHERE tablename.fieldname =
matchtablename.fieldname
```

Only very simple theta joins are supported (the **FROM** *table* clause is transformed into nested subqueries).

You can specify how **DELETE** executes by providing one or more execution options as a comma-separated list. You provide these options in a comment with the following specific syntax:

```
/* IRIS_DELETE_HINT: option,option2 */
```

Where *option* can be the following: %NOCHECK, %NOFPLAN, %NOINDEX, %NOLOCK, %NOTRIGGER, %PROFILE, %PROFILE_ALL. Refer to the InterSystems SQL [DELETE command](#) for details.

You can provide optimization hints to the **DELETE FROM** clause as a comma-separated list. You provide these hints in a comment with the following specific syntax:

```
/* IRIS_DELETEFROM_HINT: hint,hint2 */
```

Where *hint* can be the following: %ALLINDEX, %FIRSTTABLE *tablename*, %FULL, %INORDER, %IGNOREINDICES, %NOFLATTEN, %NOMERGE, %NOSVSO, %NOTOPOPT, %NOUNIONOROPT, and %STARTTABLE. Refer to the InterSystems SQL [FROM](#) clause for details.

The following *table_hints* are parsed but ignored: FASTFIRSTROW, HOLDINDEX, INDEX(name), NOLOCK, PAGLOCK, READCOMMITTED, READPAST, READUNCOMMITTED, REPEATABLEREAD, ROWLOCK, SERIALIZABLE, SHARED, TABLOCK, TABLOCKX, UPDLOCK, XLOCK. Table hints can be optionally preceded by the WITH keyword, and, if WITH is specified, optionally enclosed in parentheses. A list of table hints can be separated by either commas or blank spaces.

DELETE sets the @@ROWCOUNT system variable to the number of rows deleted, and the @@IDENTITY system variable to the IDENTITY value of the last row deleted.

The following options are not supported:

- MSSQL rowset functions.
- MSSQL OPTION clause.

4.2.2 INSERT

Inserts rows of data into a table. The following syntactic forms are supported:

```
INSERT [INTO] tablename (fieldname[,fieldname2[,...]]) VALUES (list_of_values)

INSERT [INTO] tablename (fieldname[,fieldname2[,...]]) SELECT select_list
```

The INTO keyword is optional. Specify *tablename* as described in [Table References](#).

For the VALUES syntax, the VALUES keyword is mandatory for both MSSQL and Sybase. The (*fieldname*) list is optional if the *list_of_values* lists all user-specified fields in the order defined in the table. If field names are specified, the *list_of_values* is a comma-separated list of values that matches the list of field names in number and data type.

You can specify how **INSERT** executes by providing one or more execution options as a comma-separated list. You provide these options in a comment with the following specific syntax:


```
/* IRIS_INSERT_HINT: option,option2 */
```

Where *option* can be the following: %NOCHECK, %NOFPLAN, %NOINDEX, %NOLOCK, %NOTRIGGER, %PROFILE, %PROFILE_ALL. Refer to the InterSystems SQL [INSERT command](#) for details.

The following *table_hints* are parsed but ignored: FASTFIRSTROW, HOLDINDEX, INDEX(name), NOLOCK, PAGLOCK, READCOMMITTED, READPAST, READUNCOMMITTED, REPEATABLE_READ, ROWLOCK, SERIALIZABLE, SHARED, TABLOCK, TABLOCKX, UPDLOCK, XLOCK. Table hints can be optionally preceded by the WITH keyword, and, if WITH is specified, optionally enclosed in parentheses. A list of table hints can be separated by either commas or blank spaces.

INSERT sets the @@ROWCOUNT system variable to the number of rows inserted, and the @@IDENTITY system variable to the IDENTITY value of the last row inserted.

The following options are not supported:

- (*fieldname*) DEFAULT VALUES or (*fieldname*) VALUES (DEFAULT). A field's default value is used when the field is not specified in the **INSERT** statement.
- (*fieldname*) EXECUTE procname.
- Sybase insert load option clauses: LIMIT, NOTIFY, SKIP, or START ROW ID.
- Sybase insert select load option clauses: WORD SKIP, IGNORE CONSTRAINT, MESSAGE LOG, or LOG DELIMITED BY.
- Sybase LOCATION clause.
- MSSQL INSERT TOP clause.
- MSSQL rowset functions.

4.2.3 UPDATE

Updates values of existing rows of data in a table.

```
UPDATE tablename SET fieldname=value [,fieldname2=value2[,...]]
    [FROM tablename [,tablename2]] WHERE fieldname=value

UPDATE tablename SET fieldname=value[,fieldname2=value2[,...]]
    WHERE [tablename.]fieldname=value
```

These syntactic forms are vendor-specific:

- Sybase: the optional FROM keyword syntax is used to specify an optional table (or joined tables) used in a condition. Only very simple theta joins are supported (the FROM table clause is transformed into nested subqueries).
- MSSQL: the *tablename.fieldname* syntax is used to specify an optional table used in a condition.

The *value* data type and length must match the *fieldname* defined data type and length. A *value* can be an expression that resolves to a literal value or it can be the NULL keyword. It cannot be the DEFAULT keyword.

Specify *tablename* as described in [Table References](#).

UPDATE supports the use of a local variable on the left-hand-side of a SET clause. This local variable can be either instead of a field name or in addition to a field name. The following example shows a SET to a field name, a SET to a local variable, and a SET to both a field name and a local variable:

```
UPDATE table SET x=3,@v=b,@c=Count=Count+1
```

You can specify how **UPDATE** executes by providing one or more execution options as a comma-separated list. You provide these options in a comment with the following specific syntax:

```
/* IRIS_UPDATE_HINT: option,option2 */
```

Where *option* can be the following: %NOCHECK, %NOFPLAN, %NOINDEX, %NOLOCK, %NOTRIGGER, %PROFILE, %PROFILE_ALL. Refer to the InterSystems SQL [UPDATE command](#) for details.

You can provide optimization hints to the **UPDATE FROM** clause as a comma-separated list. You provide these hints in a comment with the following specific syntax:

```
/* IRIS_UPDATEFROM_HINT: hint,hint2 */
```

Where *hint* can be the following: %ALLINDEX, %FIRSTTABLE *tablename*, %FULL, %INORDER, %IGNOREINDICES, %NOFLATTEN, %NOMERGE, %NOSVSO, %NOTOPOPT, %NOUNIONOROPT, and %STARTTABLE. Refer to the InterSystems SQL [FROM clause](#) for details.

The following *table hints* are parsed but ignored: FASTFIRSTROW, HOLDINDEX, INDEX(name), NOLOCK, PAGLOCK, READCOMMITTED, READPAST, READUNCOMMITTED, REPEATABLEREAD, ROWLOCK, SERIALIZABLE, SHARED, TABLOCK, TABLOCKX, UPDLOCK, XLOCK. Table hints can be optionally preceded by the WITH keyword, and, if WITH is specified, optionally enclosed in parentheses. A list of table hints can be separated by either commas or blank spaces.

UPDATE sets the @@ROWCOUNT system variable to the number of rows updated, and the @@IDENTITY system variable to the IDENTITY value of the last row updated.

The following Dynamic SQL example shows a simple **UPDATE** operation:

```
SET sql=9
SET sql(1)="CREATE TABLE #mytest (MyId INT PRIMARY KEY,"
SET sql(2)="LastName VARCHAR(20) CONSTRAINT nonull_lname NOT NULL,"
SET sql(3)="FirstName VARCHAR(20) DEFAULT '***TBD***') "
SET sql(4)="INSERT INTO #mytest(MyId,LastName,FirstName) VALUES (1224,'Smith','John') "
SET sql(5)="INSERT INTO #mytest(MyId,LastName) VALUES (1225,'Jones') "
SET sql(6)="INSERT INTO #mytest(MyId,LastName) VALUES (1226,'Brown') "
SET sql(7)="UPDATE #mytest SET FirstName='Fred' WHERE #mytest.LastName='Jones' "
SET sql(8)="SELECT FirstName,LastName FROM #mytest ORDER BY LastName"
SET sql(9)="DROP TABLE #mytest"
SET statement=##class(%SQL.Statement).%New()
SET statement.%Dialect="MSSQL"
SET status=statement.%Prepare(.sql)
WRITE status,!
SET result=statement.%Execute()
DO result.%Display()
```

The following options are not supported:

- Sybase ORDER BY clause.
- MSSQL OPTION clause.
- MSSQL TOP clause.
- MSSQL rowset functions.

4.2.4 READTEXT

Reads data from a stream field.

```
READTEXT tablename.fieldname textptr offset size
```

The MSSQL **READTEXT** statement returns stream data from a field of a table. It requires a valid text pointer value, which can be retrieved using the **TEXTPTR** function, as shown in the following example:

```
DECLARE @ptrval binary(16);
SELECT @ptrval = TEXTPTR(Notes) FROM Sample.Person
READTEXT Sample.Person.Notes @ptrval 0 0
```

The *textptr* must be declared as binary. A *textptr* is only defined for a text field that is not null. You can specify an initial non-null value for a text field using the [INSERT](#) statement.

The *offset* can be 0, a positive integer value, or NULL: 0 reads from the beginning of the text. A positive integer reads from the *offset* position. NULL reads from the end of the text; that is, it completes successfully but returns no value.

The *size* can be 0 or a positive integer value, or NULL: 0 reads all characters from the *offset* position to the end of the text. A positive integer reads the *size* number of characters from the *offset* position. NULL completes successfully but returns no value.

The MSSQL HOLDLOCK keyword is parsed but ignored.

4.2.5 WRITETEXT

Writes data to a stream field, replacing the existing data value.

```
WRITETEXT tablename.fieldname textptr value
```

The MSSQL **WRITETEXT** statement writes data to a stream field of a table. It requires a valid text pointer value, which can be retrieved using the **TEXTPTR** function, as shown in the following example:

```
DECLARE @ptrval binary(16);
SELECT @ptrval = TEXTPTR(Notes) FROM Sample.Person
WRITETEXT Sample.Person.Notes @ptrval 'This is the new text value'
```

The *textptr* must be declared as binary. A *textptr* is only defined for a text field that is not null. You can specify an initial non-null value for a text field using the [INSERT](#) statement.

The MSSQL BULK keyword is not supported.

The MSSQL WITH LOG keyword phrase is parsed but ignored.

4.2.6 UPDATETEXT

Updates data in a stream field.

```
UPDATETEXT tablename.fieldname textptr offset deletelength value
```

The MSSQL **UPDATETEXT** statement updates stream data from a field of a table. It requires a valid text pointer value, which can be retrieved using the **TEXTPTR** function. The following example updates the contents of the Notes stream data field by inserting the word 'New' at the beginning of the existing data value:

```
DECLARE @ptrval binary(16);
SELECT @ptrval = TEXTPTR(Notes) FROM Sample.Person
WRITETEXT Sample.Person.Notes @ptrval 0 0 'New'
```

The *textptr* must be declared as binary. A *textptr* is only defined for a text field that is not null. You can specify an initial non-null value for a text field using the [INSERT](#) statement.

The *offset* can be an integer value or NULL: 0 inserts the *value* at the beginning of the existing text. NULL inserts the *value* at the end of the existing text.

The *deletelength* can be an integer value or NULL: 0 or NULL deletes no existing characters from the *offset* position before inserting the *value*. A positive integer deletes that number of existing characters from the *offset* position before inserting the *value*.

The MSSQL BULK keyword is not supported.

The MSSQL WITH LOG keyword phrase is parsed but ignored.

4.2.7 TRUNCATE TABLE

Deletes all of the data from a table.

```
TRUNCATE TABLE tablename
```

Invokes the InterSystems SQL [TRUNCATE TABLE](#) command, which deletes all rows from the specified table and resets the RowId (ID), IDENTITY, and SERIAL (%Counter) row counters and the stream field OID counter values.

You can specify how **TRUNCATE TABLE** executes by providing one or more execution options as a comma-separated list. You provide these options in a comment with the following specific syntax:

```
/* IRIS_DELETE_HINT: option,option2 */
```

Where *option* can be the following: %NOCHECK, %NOLOCK. Refer to the InterSystems SQL [TRUNCATE TABLE](#) for details.

4.3 Query Statements

4.3.1 SELECT

```
SELECT [DISTINCT | ALL]
  [TOP [( ){ int | @var | ? | ALL} [ ] ] ]
  select-item { ,select-item }
  [INTO [(#)copytable] ]
  [FROM tablename [(AS) t-alias] [, tablename2 [(AS) t-alias2] ] ]
  [[WITH] [( ) tablehint=val [, tablehint=val] [ ] ] ]
  [WHERE condition-expression]
  [GROUP BY scalar-expression]
  [HAVING condition-expression]
  [ORDER BY item-order-list [ASC | DESC] ]
```

The above SELECT syntax is supported. The following features are not supported:

- TOP *nn* PERCENT or TOP WITH TIES
- OPTION
- WITH CUBE
- WITH ROLLUP
- GROUP BY ALL
- GROUP WITH
- COMPUTE clause
- FOR BROWSE

TOP *nn* specifies the number of rows to retrieve. InterSystems TSQL supports TOP *nn* with a integer, ?, local variable, or the keyword ALL. The TOP argument can be enclosed in parentheses TOP (*nn*). These parentheses are retained, preventing preparer substitution. If SET ROWCOUNT specifies fewer rows than TOP *nn*, the SET ROWCOUNT value is used. The following Dynamic SQL example shows the use of TOP with a local variable:

```

SET sql=3
SET sql(1)="DECLARE @var INT"
SET sql(2)="SET @var=4"
SET sql(3)="SELECT TOP @var Name, Age FROM Sample.Person"
SET statement=##class(%SQL.Statement).%New()
SET statement.%Dialect="MSSQL"
SET status=statement.%Prepare(.sql)
SET result=statement.%Execute()
DO result.%Display()

```

The *select-item* list can contain the following:

- field names, functions, and expressions
- the **\$IDENTITY** pseudo-field name, which always returns the [RowID](#) value, regardless of the field name assigned to the RowID.
- an asterisk: **SELECT *** is supported. The asterisk means to select all fields in the specified table. You can qualify the asterisk with the table name or table alias: **SELECT mytable.***.
- a subquery
- stream fields. A **SELECT** on a stream field returns the *oref* (object reference) of the opened stream object.

An INTO clause can be used to copy data from an existing table into a new table. By default, **SELECT** creates the INTO table with the same field names and data types as the fields selected from the source table. The INTO table cannot already exist. This INTO table can be a permanent table, or a temporary table, as shown in the following examples:

```
SELECT Name INTO Sample.NamesA_G FROM Sample.Person WHERE name LIKE '[A-G]%'
```

```

SELECT Name INTO #MyTemp FROM Sample.Person WHERE name LIKE '[A-G]%'
SELECT * FROM #MyTemp

```

You can specify a different name for an INTO table field by using a field alias, as shown in the following example:

```
SELECT Name AS Surname INTO Sample.NamesA_G FROM Sample.Person WHERE name LIKE '[A-G]%'
```

An INTO clause cannot be used when the **SELECT** is a subquery or is part of a UNION.

The FROM clause is not required. A **SELECT** without a FROM clause can be used to assign a value to a local variable, as follows:

```

DECLARE @myvar INT
SELECT @myvar=1234
PRINT @myvar

```

The FROM clause supports table hints with either of the following syntactic forms:

```

FROM tablename (INDEX=indexname)
FROM tablename INDEX (indexname)

```

Table hints can be optionally preceded by the WITH keyword, and optionally enclosed in parentheses. A list of table hints can be separated by either commas or blank spaces. The following table hints are parsed but ignored: FASTFIRSTROW, HOLDINDEX, NOLOCK, PAGLOCK, READCOMMITTED, READPAST, READUNCOMMITTED, REPEAT-ABLEREAD, ROWLOCK, SERIALIZABLE, SHARED, TABLOCK, TABLOCKX, UPDLOCK, XLOCK.

You can provide optimization hints to the **SELECT FROM** clause as a comma-separated list. You provide these hints in a comment with the following specific syntax:

```
/* IRIS_SELECTFROM_HINT: hint,hint2 */
```

Where *hint* can be the following: %ALLINDEX, %FIRSTTABLE *tablename*, %FULL, %INORDER, %IGNOREINDICES, %NOFLATTEN, %NOMERGE, %NOSVSO, %NOTOPOPT, %NOUNIONOROPT, and %STARTTABLE. Refer to the InterSystems SQL [FROM](#) clause for details.

A WHERE clause can use AND, OR, and NOT logic keywords. It can group multiple search conditions using parentheses. The WHERE clause supports the following search conditions:

- Equality comparisons: = (equals), <> (not equals), < (less than), > (greater than), <= (less than or equals), >= (greater than or equals).
- IS NULL and IS NOT NULL comparisons.
- BETWEEN comparisons: Age BETWEEN 21 AND 65 (inclusive of 21 and 65); Age NOT BETWEEN 21 AND 65 (exclusive of 21 and 65). BETWEEN is commonly used for a range of numeric values, which collate in numeric order. However, BETWEEN can be used for a collation sequence range of values of any data type. It uses the same collation type as the field it is matching against. By default, string data types collate as not case-sensitive.
- IN comparisons: Home_State IN ('MA', 'RI', 'CT').
- LIKE and NOT LIKE comparisons, specified as a quoted string. The comparison string can contain wildcards: _ (any single character); % (any string); [abc] (any value in the set specified as a list of items); [a-c] (any value in the set specified as a range of items). InterSystems TSQL does not support the ^ wildcard. A LIKE comparison can include an ESCAPE clause, such as the following: WHERE CategoryName NOT LIKE 'D_%' ESCAPE '\'.
- EXISTS comparison check: used with a subquery to test whether the subquery evaluates to the empty set. For example SELECT Name FROM Sample.Person WHERE EXISTS (SELECT LastName FROM Sample.Employee WHERE LastName='Smith'). In this example, all Names are returned from Sample.Person if a record with LastName='Smith' exists in Sample.Employee. Otherwise, no records are returned from Sample.Person.
- ANY and ALL comparison check: used with a subquery and an equality comparison operator. The SOME keyword is a synonym for ANY.

WHERE clause and HAVING clause comparisons are not case-sensitive.

A HAVING clause can be specified after a GROUP BY clause. The HAVING clause is like a WHERE clause that can operate on groups, rather than on the full data set. HAVING and WHERE use the same comparisons. This is shown in the following example:

```
SELECT Home_State, MIN(Age) AS Youngest,
       AVG(Age) AS AvgAge, MAX(Age) AS Oldest
FROM Sample.Person
GROUP BY Home_State
HAVING Age < 21
ORDER BY Youngest
```

The following Dynamic SQL example selects table data into a result set:

```
SET sql=7
SET sql(1)="CREATE TABLE #mytest (MyId INT PRIMARY KEY, "
SET sql(2)="LastName VARCHAR(20), "
SET sql(3)="FirstName VARCHAR(20))"
SET sql(4)="INSERT INTO #mytest(MyId,LastName,FirstName) VALUES (1224,'Smith','John') "
SET sql(5)="INSERT INTO #mytest(MyId,LastName,FirstName) VALUES (1225,'Jones','Wilber') "
SET sql(6)="SELECT FirstName,LastName FROM #mytest"
SET sql(7)="DROP TABLE #mytest"
SET statement=##class(%SQL.Statement).%New()
SET statement.%Dialect="MSSQL"
SET status=statement.%Prepare(.sql)
SET result=statement.%Execute()
DO result.%Display()
```

The following Dynamic SQL example selects a single field value into a local variable:

```

SET sql=9
SET sql(1)="CREATE TABLE #mytest (MyId INT PRIMARY KEY,"
SET sql(2)="LastName VARCHAR(20),"
SET sql(3)="FirstName VARCHAR(20))"
SET sql(4)="INSERT INTO #mytest(MyId,LastName,FirstName) VALUES (1224,'Smith','John')"
```

```

SET sql(5)="INSERT INTO #mytest(MyId,LastName,FirstName) VALUES (1225,'Jones','Wilber')"
```

```

SET sql(6)="DECLARE @nam VARCHAR(20)"
SET sql(7)="SELECT @nam=LastName FROM #mytest"
SET sql(8)="PRINT @nam"
SET sql(9)="DROP TABLE #mytest"
SET statement=##class(%SQL.Statement).%New()
SET statement.%Dialect="MSSQL"
SET status=statement.%Prepare(.sql)
DO statement.%Execute()
```

An ORDER BY clause can specify ascending (ASC) or descending (DESC) order. The default is ascending. Unlike InterSystems SQL, an ORDER BY may be used in subqueries and in queries that appear in expressions. For example:

```
SET @var = (SELECT TOP 1 name FROM mytable ORDER BY name)
```

4.3.2 JOIN

JOIN (equivalent to INNER JOIN), INNER JOIN, and LEFT JOIN supported. Parentheses can be used to rationalize parsing of multiple joins.

4.3.3 UNION

A union of two (or more) **SELECT** statements is supported. InterSystems TSQL supports **UNION** and **UNION ALL**. If you specify **UNION ALL**, only the first **SELECT** can specify an INTO table. This INTO table can be a defined table, or a temporary table generated from the **SELECT** field list.

4.3.4 FETCH Cursor

The **OPEN**, **FETCH**, **CLOSE**, and **DEALLOCATE** commands are mainly supported. The following features are not supported:

- OPEN/FETCH/CLOSE @local
- FETCH followed by any qualifier other than NEXT (the qualifier can be omitted).
- Note that DEALLOCATE is supported, but that, by design, it generates no code.

4.4 Flow of Control Statements

4.4.1 IF

Executes a block of code if a condition is true.

The **IF** command is supported with four syntactic forms:

IF...ELSE syntax:

```
IF condition
statement
[ELSE statement]
```

IF...THEN...ELSE single-line syntax:

```
IF condition THEN statement [ELSE statement]
```

ELSEIF...END IF syntax:

```
IF condition THEN
statements
{ELSEIF condition THEN statements}
[ELSE statements]
END IF
```

ELSE IF (SQL Anywhere) syntax:

```
IF condition THEN statement
{ELSE IF condition THEN statement}
[ELSE statement]
```

The first syntactic form is the TSQL standard format. No THEN keyword is used. You may use white space and line breaks freely. To specify more than one *statement* in a clause you must use BEGIN and END keywords to demarcate the block of statements. The ELSE clause is optional. This syntax is shown in the following example:

```
SET sql=4
SET sql(1)="DECLARE @var INT"
SET sql(2)="SET @var=RAND()"
SET sql(3)="IF @var<.5 PRINT 'The Oracle says No'"
SET sql(4)="ELSE PRINT 'The Oracle says Yes' "
SET statement=##class(%SQL.Statement).%New()
SET statement.%Dialect="MSSQL"
SET status=statement.%Prepare(.sql)
SET result=statement.%Execute()
DO result.%Display()
```

The second syntactic form is single-line syntax. The THEN keyword is required. A line break restriction requires that IF condition THEN statement all be on the same line, though only the first keyword of the *statement* must be on that line. Otherwise, you may use white space and line breaks freely. To specify more than one *statement* in a clause you must use BEGIN and END keywords to demarcate the block of statements. The ELSE clause is optional. This syntax is shown in the following example:

```
SET sql=3
SET sql(1)="DECLARE @var INT "
SET sql(2)="SET @var=RAND() "
SET sql(3)="IF @var<.5 THEN PRINT 'No' ELSE PRINT 'Yes' "
SET statement=##class(%SQL.Statement).%New()
SET statement.%Dialect="MSSQL"
SET status=statement.%Prepare(.sql)
SET result=statement.%Execute()
DO result.%Display()
```

The third syntactic form provides an ELSEIF clause. You can specify zero, one, or more than one ELSEIF clauses, each with its own *condition* test. Within an IF, ELSEIF, or ELSE clause you can specify multiple statements. BEGIN and END keywords are permitted but not required. A line break restriction requires a line break between IF condition THEN and the first *statement*. Otherwise, you may use white space and line breaks freely. The ELSE clause is optional. The END IF keyword clause is required. This syntax is shown in the following example:


```

SET sql=14
SET sql(1)="DECLARE @var INT "
SET sql(2)="SET @var=RAND() "
SET sql(3)="IF @var<.2 THEN "
SET sql(4)="PRINT 'The Oracle' "
SET sql(5)="PRINT 'says No' "
SET sql(6)="ELSEIF @var<.4 THEN "
SET sql(7)="PRINT 'The Oracle' "
SET sql(8)="PRINT 'says Possibly' "
SET sql(9)="ELSEIF @var<.6 THEN "
SET sql(10)="PRINT 'The Oracle' "
SET sql(11)="PRINT 'says Probably' "
SET sql(12)="ELSE PRINT 'The Oracle' "
SET sql(13)="PRINT 'says Yes' "
SET sql(14)="END IF"
SET statement=##class(%SQL.Statement).%New()
SET statement.%Dialect="MSSQL"
SET status=statement.%Prepare(.sql)
SET result=statement.%Execute()
DO result.%Display()

```

The fourth syntactic form is compatible with SQL Anywhere. It provides an ELSE IF clause (note space between keywords). You can specify zero, one, or more than one ELSE IF clauses, each with its own *condition* test. To specify more than one *statement* in a clause you must use BEGIN and END keywords to demarcate the block of statements. You may use white space and line breaks freely. The ELSE clause is optional. This syntax is shown in the following example:

```

SET sql=6
SET sql(1)="DECLARE @var INT "
SET sql(2)="SET @var=RAND() "
SET sql(3)="IF @var<.2 THEN PRINT 'The Oracle says No'"
SET sql(4)="ELSE IF @var<.4 THEN PRINT 'The Oracle says Possibly'"
SET sql(5)="ELSE IF @var<.6 THEN PRINT 'The Oracle says Probably'"
SET sql(6)="ELSE PRINT 'The Oracle says Yes'"
SET statement=##class(%SQL.Statement).%New()
SET statement.%Dialect="MSSQL"
SET status=statement.%Prepare(.sql)
SET result=statement.%Execute()
DO result.%Display()

```

4.4.2 WHILE

Repeatedly executes a block of code while a condition is true.

```
WHILE condition BEGIN statements END
```

The BREAK keyword exits the WHILE loop.

The CONTINUE keyword immediately returns to the top of the WHILE loop.

The BEGIN and END keywords are required if *statements* is more than one command.

The following example returns four result sets, each containing a pair of records in ascending ID sequence:

```

DECLARE @n INT;
SET @n=0;
WHILE @n<8 BEGIN
    SELECT TOP 2 ID,Name FROM Sample.Person WHERE ID>@n
    SET @n=@n+2
END;

```

4.4.3 CASE

Returns a value from the first match of multiple specified values.

```
CASE expression WHEN value THEN rtnval
[WHEN value2 THEN rtnval2] [...]
[ELSE rtndefault]
END
```

The **WHEN** *value* must be a simple value. It cannot be a boolean expression.

The **ELSE** clause is optional. If no **WHEN** clause is satisfied and the **ELSE** clause is not provided, the **CASE** statement returns *expression* as **NULL**.

For example:

```
SELECT CASE Name WHEN 'Fred Rogers' THEN 'Mr. Rogers'
              WHEN 'Fred Astaire' THEN 'Ginger Rogers'
              ELSE 'Somebody Else' END
FROM Sample.Person
```

The returned value does not have to match the data type of *expression*.

CASE parses but ignores **WHEN NULL THEN rtnval** cases.

4.4.4 GOTO and Labels

InterSystems TSQL supports the **GOTO** command and labels. A label must be a valid [TSQL identifier](#) followed by a colon (:). A **GOTO** reference to a label does not include the colon.

4.4.5 WAITFOR

Used to delay execution until a specific elapse of time or clock time.

```
WAITFOR DELAY timeperiod
WAITFOR TIME clocktime
```

timeperiod is the amount of time to wait before resuming execution, expressed as 'hh:mm[:ss[.fff]]'. Thus **WAITFOR DELAY '00:00:03'** provides a time delay of 3 seconds; **WAITFOR DELAY '00:03'** provides a time delay of 3 minutes; **WAITFOR DELAY '00:00:00.9'** provides a time delay of nine-tenths of a second. Note that the fractional second divider is a period, not a colon.

clocktime is the time at which to resume execution, expressed as 'hh:mm[:ss[.fff]]', using a 24-hour clock. Thus **WAITFOR TIME '14:35:00'** resumes execution at 2:35pm; **WAITFOR TIME '00:00:03'** resumes execution at 3 seconds after midnight.

The following options are not supported:

- Sybase **CHECK EVERY** clause.
- Sybase **AFTER MESSAGE BREAK** clause.
- MSSQL **RECEIVE** clause.

4.5 Assignment Statements

4.5.1 DECLARE

Declares the data type for a local variable.

```
DECLARE @var [AS] datatype [ = initval]
```

Only the form which declares local variables is supported; cursor variables are not supported. The AS keyword is optional. Unlike InterSystems SQL, you must declare a local variable before you can set it.

@var can be any local variable name. Local variable names are not case-sensitive.

The *datatype* can be any valid data type, such as CHAR(12) or INT. TEXT, NTEXT, and IMAGE data types are not allowed. For further details on data types, refer to the [TSQL Constructs](#) chapter of this document.

The optional *initval* argument allows you to set the initial value of the local variable. You can set it to a literal value or to any of the following: NULL, USER, CURRENT DATE (or CURRENT_DATE), CURRENT TIME (or CURRENT_TIME), CURRENT TIMESTAMP (or CURRENT_TIMESTAMP), or CURRENT_USER. The DEFAULT and CURRENT_DATABASE keywords are not supported. Alternatively, you can set the value of a local value using the [SET](#) command or the [SELECT](#) command. For example:

```
DECLARE @c INT;
SELECT @c=100;
```

You can specify multiple local variable declarations as a comma-separated list. Each declaration must have its own data type and (optionally) its own initial value:

```
DECLARE @a INT=1, @b INT=2, @c INT=3
```

4.5.2 SET

Assigns a value to a local variable or an environment setting.

Used to assign a value to a local variable:

```
DECLARE @var CHAR(20)
SET @var='hello world'
```

Used to set an environment setting:

```
SET option ON
```

These settings have immediate effect at parse time, whether inside a stored procedure or not. The change persists until another **SET** command alters it – even if the **SET** is made inside a stored procedure, and accessed outside the SP or in another SP.

The following **SET** environment settings are supported:

- **SET ANSI_NULLS** Permitted values are SET ANSI_NULLS ON and SET ANSI_NULLS OFF. If ANSI_NULLS OFF, $a=b$ is true if ($a=b$ OR (a IS NULL) AND (b IS NULL)). See the [ANSI_NULLS](#) TSQL system-wide configuration setting.
- **SET DATEFIRST integer** specifies which day is treated as the first day of the week. Permitted values are 1 through 7, with 1=Monday and 7=Sunday. The default is 7.

- `SET IDENTITY_INSERT` Permitted values are `SET IDENTITY_INSERT ON` and `SET IDENTITY_INSERT OFF`. If ON, an INSERT statement can specify an identity field value. This variable applies exclusively to the current process and cannot be set on linked tables. Therefore, to use this option you should define a procedure in TSQL to perform both the `SET IDENTITY_INSERT` and the INSERT, then link the procedure and execute the procedure in InterSystems IRIS via the gateway.
- `SET NOCOUNT` Permitted values are `SET NOCOUNT ON` and `SET NOCOUNT OFF`. When set to ON, messages indicating the number of rows affected by a query are suppressed. This can have significant performance benefits.
- `SET QUOTED_IDENTIFIER` Permitted values are `SET QUOTED_IDENTIFIER ON` and `SET QUOTED_IDENTIFIER OFF`. When `SET QUOTED_IDENTIFIER` is on, double quotes are parsed as delimiting a quoted identifier. When `SET QUOTED_IDENTIFIER` is off, double quotes are parsed as delimiting a string literal. The preferable delimiters for string literals are single quotes. See the [QUOTED_IDENTIFIER](#) TSQL system-wide configuration setting.
- `SET ROWCOUNT` Set to an integer. Affects subsequent SELECT, INSERT, UPDATE, or DELETE statements to limit the number of rows affected. In a SELECT statement, ROWCOUNT takes precedence over TOP: if ROWCOUNT is less than TOP, the ROWCOUNT number of rows is returned; if TOP is less than ROWCOUNT, the TOP number of rows is returned. ROWCOUNT remains set for the duration of the process or until you revert it to default behavior. To revert to default behavior, `SET ROWCOUNT 0`. If you specify a fractional value, ROWCOUNT is set to the next larger integer.
- `SET TRANSACTION ISOLATION LEVEL` See [Transaction Statements](#) below.

The following **SET** environment setting is parsed, but ignored:

- `SET TEXTSIZE integer`

4.6 Transaction Statements

InterSystems TSQL provides support for transactions, including named transaction names. It does not support savepoints. Distributed transactions are not supported.

4.6.1 SET TRANSACTION ISOLATION LEVEL

Supported for the following forms only:

- `SET TRANSACTION ISOLATION LEVEL READ COMMITTED`
- `SET TRANSACTION ISOLATION LEVEL READ UNCOMMITTED.`

READ VERIFIED and other options are not supported.

Sybase `SET TRANSACTION ISOLATION LEVEL n` integer option codes (0, 1, 2, 3) are not supported.

4.6.2 BEGIN TRANSACTION

Begins the current transaction.

```
BEGIN TRAN [name]
BEGIN TRANSACTION [name]
```

Initiates a transaction. The optional *name* argument can be used to specify a named transaction, also known as a savepoint. The *name* value must be supplied as a literal; it cannot be a variable.

You can issue multiple **BEGIN TRANSACTION** statements to create multiple nested transactions. You can use the @@trancount special variable to determine the current transaction level. Each transaction level must be resolved by a **COMMIT** statement or a **ROLLBACK** statement.

Note: A [Data Management Language \(DML\) statement](#) that is within an explicit transaction *cannot* resolve an unqualified table name using a schema search path.

4.6.3 COMMIT TRANSACTION

Commits the current transaction.

```
COMMIT
COMMIT TRAN
COMMIT TRANSACTION
COMMIT WORK
```

These four syntactical forms are functionally identical; the **COMMIT** keyword, as specified below, refers to any of these syntactical forms. A **COMMIT** statement commits all work completed during the current transaction, resets the transaction level counter, and releases all locks established. This completes the transaction. Work committed cannot be rolled back.

If multiple **BEGIN TRANSACTION** statements have created nested transactions, **COMMIT** completes the current nested transaction. A transaction is defined as the operations since and including the **BEGIN TRANSACTION** statement. A **COMMIT** restores the transaction level counter to its state immediately prior to the **BEGIN TRANSACTION** statement that initialized the transaction. You can use the @@trancount special variable to determine the current transaction level.

A **COMMIT** cannot specify a named transaction. If you specify a transaction name as part of a **COMMIT** statement, the presence of this name is parsed without issuing an error, but the transaction name is not validated and it is ignored.

Sybase performs no operation and does not issue an error if a **COMMIT** is issued when not in a transaction.

4.6.4 ROLLBACK TRANSACTION

Rolls back the specified transaction or all current transactions.

```
ROLLBACK [name]
ROLLBACK TRAN [name]
ROLLBACK TRANSACTION [name]
ROLLBACK WORK [name]
```

These four syntactical forms are functionally identical; the **ROLLBACK** keyword, as specified below, refers to any of these syntactical forms. The optional *name* argument specifies a named transaction, as specified by a **BEGIN TRANSACTION name** statement. The *name* value must be supplied as a literal; it cannot be a variable.

A **ROLLBACK** rolls back a transaction, undoing work performed but not committed, decrementing the transaction level counter, and releasing locks. It is used to restore the database to a previous consistent state.

- A **ROLLBACK** rolls back all work completed during the current transaction (or series of nested transactions), resets the transaction level counter to zero and releases all locks. This restores the database to its state before the beginning of the outermost nested transaction.
- A **ROLLBACK name** rolls back all work done since the specified named transaction (savepoint) and decrements the transaction level counter by the number of savepoints undone. When all savepoints have been either rolled back or committed and the transaction level counter reset to zero, the transaction is completed. If the named transaction does not exist, or has already been rolled back, **ROLLBACK** rolls back the entire current transaction.

Sybase performs no operation and does not issue an error if a **ROLLBACK** is issued when not in a transaction.

4.6.5 SAVE TRANSACTION

The **SAVE TRANSACTION** [*savepoint-name*] statement is parsed but ignored in InterSystems TSQL. It performs no operation.

4.6.6 LOCK TABLE

Enables the current user to lock a table.

```
LOCK TABLE tablename IN {SHARE | EXCLUSIVE} MODE [WAIT numsecs | NOWAIT]
```

The **LOCK TABLE** statement locks all of the records in the specified table. You can lock a table in SHARE MODE or in EXCLUSIVE MODE. The optional WAIT clause specifies the number of seconds to wait in attempting to acquire the table lock. The **LOCK TABLE** statement immediately releases any prior lock held by the current user on the specified table.

LOCK TABLE is only meaningful within a transaction. It locks the table for the duration of the current transaction. When not in a transaction, **LOCK TABLE** performs no operation.

Specify *tablename* as described in [Table References](#). **LOCK TABLE** supports locking a single table; it does not support locking multiple tables.

LOCK TABLE supports SHARE and EXCLUSIVE modes; it does not support WRITE mode.

LOCK TABLE does not support the WITH HOLD clause.

WAIT time is specified as an integer number of seconds; **LOCK TABLE** does not support WAIT time specified as clock time.

4.7 Procedure Statements

The following standard Transact-SQL statements are supported.

4.7.1 CREATE PROCEDURE / CREATE FUNCTION

Creates a named executable procedure.

```
CREATE PROCEDURE procname [[@var [AS] datatype [= | DEFAULT value] [,...]] [RETURNS datatype] [AS] code
CREATE PROC procname [[@var [AS] datatype [= | DEFAULT value] [,...]] [RETURNS datatype] [AS] code
CREATE FUNCTION procname [[@var [AS] datatype [= | DEFAULT value] [,...]] [RETURNS datatype] [AS] code
```

You can return a single scalar value result from either a PROCEDURE or a FUNCTION. OUTPUT parameters and default values are also supported. These commands convert the return type from a TSQL type declaration to an InterSystems IRIS type descriptor. Currently, result sets and tables can't be returned.

Supported as either **CREATE PROCEDURE** or **CREATE PROC**. **CREATE FUNCTION** is very similar to **CREATE PROCEDURE**, but the routine type argument value is "FUNCTION", rather than "PROCEDURE".

- Any statements can be used in a **CREATE FUNCTION**.
- The RETURN keyword is allowed in a **CREATE PROCEDURE**. If a procedure completes without invoking a **RETURN** or **RAISERROR** statement, it returns an integer value of 0.

- The WITH EXECUTE keyword clause is allowed in a **CREATE PROCEDURE** and **CREATE FUNCTION**. It must appear after the RETURN keyword.

A **CREATE PROCEDURE** can specify a formal parameter list. Formal parameters are specified as a comma-separated list. Enclosing parentheses are optional. The AS keyword between the parameter variable and its data type is optional. Optionally, you can use the DEFAULT keyword or = symbol to assign a default value to a formal parameter; if no actual parameter value is specified, this default value is used. In TSQL an input formal parameter has no keyword indicator; an output formal parameter can be specified by the OUTPUT keyword following the data type. Alternatively, these formal parameters can be prefaced by the optional keywords IN, OUT, or INOUT.

The following example shows the creation of the procedure AvgAge with two formal parameters:

```
CREATE PROCEDURE AvgAge @min INT, @max INT
AS
BEGIN TRY
    SELECT AVG(Age) FROM Sample.Person
    WHERE Age > @min AND Age < @max
END TRY
BEGIN CATCH
    PRINT 'error!'
END CATCH
```

The following statement executes this procedure. In this case, the specified actual parameter values limit the averaging to ages 21 through 65:

```
EXEC AvgAge 20,66
```

The following example creates a procedure that returns the results of a division operation. The RETURNS keyword limits the number of decimal digits in the return value:

```
CREATE PROCEDURE SQLUser.MyDivide @a INTEGER, @b INTEGER, OUT @rtn INTEGER RETURNS DECIMAL(2,3)
BEGIN
SET @rtn = @a / @b;
RETURN @rtn;
END
```

The following statement executes this procedure:

```
SELECT SQLUser.MyDivide(7,3)
```

The following example shows the creation of procedure OurReply:

```
CREATE PROCEDURE OurReply @var CHAR(16) DEFAULT 'No thanks' AS PRINT @var
```

When executed without a parameter, OurReply prints the default text (“No thanks”); when executed with a parameter OurReply prints the actual parameter value specified in the **EXEC** statement.

Note that **CREATE FUNCTION** and **CREATE PROCEDURE** cannot be issued from a stored procedure.

4.7.1.1 Importing a CREATE PROCEDURE

If imported TSQL source contains a CREATE PROC statement, then a class method containing the CREATE PROC source will be created. This class method is either placed in an existing class, or in a new class whose name is based on the schema and procedure name.

If the procedure already exists, the existing implementation is replaced. If a class matching the class name generated from the schema and procedure already exists, it is used if it was previously generated by the TSQL utility. If not, then a unique class name is generated, based on the schema and procedure name. The schema defaults to the default schema defined in the system configuration. The resulting class is compiled once the procedure has been successfully created.

If logging is requested, the source statements are logged along with the name of the containing class, class method, and the formal arguments generated. Any errors encountered by the process are also reported in the log. If errors are detected during CREATE PROC processing and a new class was generated, that class is deleted.

4.7.2 ALTER FUNCTION

Supported. The WITH EXECUTE keyword clause is supported.

4.7.3 DROP FUNCTION

Deletes a function or a comma-separated list of functions.

```
DROP FUNCTION funcname [,funcname2 [,...]]
```

The IF EXISTS clause is not supported.

4.7.4 DROP PROCEDURE

Deletes a procedure or a comma-separated list of procedures.

```
DROP PROCEDURE [IF EXISTS] procname [,procname2 [,...]]
DROP PROC [IF EXISTS] procname [,procname2 [,...]]
```

The optional IF EXISTS clause suppresses errors if you specify a non-existent *procname*. If this clause is not specified, an SQLCODE -362 error is generated if you specify a non-existent *procname*. **DROP PROCEDURE** is an atomic operation; either all specified procedures are successfully deleted or none are deleted.

4.7.5 RETURN

Halts execution of a query or procedure. Can be argumentless or with an argument. Argumentless **RETURN** must be used when exiting a TRY or CATCH block. When returning from a procedure, **RETURN** can optionally return an integer status code. If you specify no status code, it returns the empty string ('').

4.7.6 EXECUTE

Executes a procedure, or executes a string of TSQL commands.

```
EXECUTE [@rtnval = ] procname [param1 [,param2 [,...]]]
EXECUTE ( 'TSQL_commands' )
```

EXEC is a synonym for EXECUTE.

- **EXECUTE *procname*** can be used to execute a stored procedure. Parameters are supplied as a comma-separated list. This parameter list is not enclosed in parentheses. Named parameters are supported.

EXECUTE *procname* can optionally receive a RETURN value, using the EXECUTE @rtn=Sample.MyProc param1,param2 syntax.

EXECUTE *procname* is similar to the [CALL](#) statement, which can also be used to execute a stored procedure. CALL uses an entirely different syntax.

```
CREATE PROCEDURE Sample.AvgAge @min INT, @max INT
AS
SELECT Name, Age, AVG(Age) FROM Sample.Person
WHERE Age > @min AND Age < @max
RETURN 99
```

```
DECLARE @rtn INT;
EXECUTE @rtn=Sample.AvgAge 18,65
SELECT @rtn
```


If the specified procedure does not exist, an SQLCODE -428 error (Stored procedure not found) is issued.

The WITH RECOMPILE clause is parsed, but ignored.

The following **EXECUTE *procname*** features are not supported: procedure variables, and procedure numbers (i.e. 'n').

- **EXECUTE (TSQL commands)** can be used to execute dynamic SQL. The TSQL command(s) are enclosed in parentheses. The TSQL commands to be executed are specified as a string enclosed in single quote characters. A TSQL command string can contain line breaks and white space. Dynamic TSQL runs in the current context.

```
EXECUTE('SELECT TOP 4 Name, Age FROM Sample.Person')
```

or

```
DECLARE @DynTopSample VARCHAR(200)
SET @DynTopSample='SELECT TOP 4 Name, Age FROM Sample.Person'
EXECUTE (@DynTopSample)
```

The following example shows an EXECUTE that returns multiple result sets:

```
EXECUTE('SELECT TOP 4 Name FROM Sample.Person
        SELECT TOP 6 Age FROM Sample.Person')
```

4.7.7 CALL

Executes a procedure.

```
[@var = ] CALL procname ([param1 [, param2 [, ...] ] ])
```

The **CALL** statement is functionally identical to the **EXECUTE *procname*** statement. It differs syntactically.

The procedure parameters are optional. The enclosing parentheses are mandatory.

The optional *@var* variable receives the value returned by the RETURN statement. If execution of the stored procedure does not conclude with a RETURN statement, *@var* is set to 0.

The following example calls a stored procedure, passing two input parameters. It receives a value from the procedure's RETURN statement:

```
DECLARE @rtn INT
@rtn=CALL Sample.AvgAge(18, 34)
SELECT @rtn
```

4.8 Other Statements

4.8.1 CREATE USER

CREATE USER creates a new user.

```
CREATE USER username
```

Executing this statement creates an InterSystems IRIS user with its password set to the specified user name. You can then use the Management Portal **System Administration** interface to change the password. You cannot explicitly set a password using CREATE USER.

User names are not case-sensitive. InterSystems TSQL and InterSystems SQL both use the same set of defined user names. InterSystems IRIS issues an error message if you try to create a user that already exists.

By default, a user has no privileges. Use the GRANT command to give privileges to a user.

4.8.2 GRANT

Grants privileges to a user or list of users.

```
GRANT privelist ON tablelist TO granteelist
```

```
GRANT EXECUTE ON proclist TO granteelist
```

- *privelist*: a single privilege or a comma-separated list of privileges. The available privileges are SELECT, INSERT, DELETE, UPDATE, REFERENCES, and ALL PRIVILEGES. ALL is a synonym for ALL PRIVILEGES. The ALTER privilege is not supported directly, but is one of the privileges granted by ALL PRIVILEGES.
- *tablelist*: a single table name (or view name) or a comma-separated list of table names and view names. Specify a table name as described in [Table References](#).
- *proclist*: a single SQL procedure or a comma-separated list of SQL procedures. All listed procedures must exist, otherwise an SQLCODE -428 error is returned.
- *granteelist*: a single grantee (user to be assigned privileges) or a comma-separated list of grantees. A grantee can be a user name, "PUBLIC" or "*". Specifying * grants the specified privileges to all existing users. A user created using CREATE USER initially has no privileges. Specifying a non-existent user in a comma-separated list of grantees has no effect; GRANT ignore that user and grants the specified privileges to the existing users in the list.

Specifying privileges for specified fields is not supported.

The WITH GRANT OPTION clause is parsed but ignored.

Granting a privilege to a user that already has that privilege has no effect and no error is issued.

4.8.3 REVOKE

Revokes granted privileges from a user or list of users.

```
REVOKE privelist ON tablelist FROM granteelist CASCADE
```

```
REVOKE EXECUTE ON proclist FROM granteelist
```

Revoking a privilege from a user that does not have that privilege has no effect and no error is issued.

See [GRANT](#) for further details.

4.8.4 PRINT

Displays the specified text to the current device.

```
PRINT expression [,expression2 [,...]]
```

An *expression* can be a literal string enclosed in single quotes, a number, or a variable or expression that resolves to a string or a number. You can specify any number of comma-separated expressions.

PRINT does not support the Sybase *arg-list* syntax. A placeholder such as %3! in an *expression* string is not substituted for, but is displayed as a literal.

4.8.5 RAISERROR

```
RAISERROR errnum 'message'
RAISERROR(error,severity,state,arg) WITH LOG
```

Both syntactic forms (with and without parentheses) are supported. Both spellings, **RAISERROR** and **RAISEERROR**, are supported and synonymous. **RAISERROR** sets the value of @@ERROR to the specified error number and error message and invokes the **%SYSTEM.Error.FromXSQL()** method.

The Sybase-compatible syntax (without parentheses) requires an *errnum* error number, the other arguments are optional.

```
RAISERROR 123 'this is a big error'
PRINT @@ERROR
```

A **RAISERROR** command raises an error condition; it is left to the user code to detect this error. However, if **RAISERROR** appears in the body of a TRY block, it transfers control to the paired CATCH block. If **RAISERROR** appears in a CATCH block it transfers control either to an outer CATCH block (if it exists) or to the procedure exit. **RAISERROR** does not trigger an exception outside of the procedure. It is up to the caller to check for the error.

When an AFTER statement level trigger executes a **RAISEERROR**, the returned **%msg** value contains the *errnum* and *message* values as message string components separated by a comma: **%msg="errnum,message"**.

The Microsoft-compatible syntax (with parentheses) requires an *error* (either an error number or a quoted error message). If you do not specify an error number, it defaults to 50000. The optional *severity* and *state* arguments take integer values.

```
RAISERROR('this is a big error',4,1) WITH LOG
PRINT @@ERROR
```

4.8.6 UPDATE STATISTICS

Optimizes query access for a specified table. The specified table can be a standard table or a # temporary table (see **CREATE TABLE** for details.) InterSystems IRIS passes the specified table name argument to the **\$SYSTEM.SQL.TuneTable()** method for optimization. **UPDATE STATISTICS** calls **\$SYSTEM.SQL.TuneTable()** with *update=1* and *display=0*. The returned **%msg** is ignored and *KeepClassUpToDate* defaults to 'false'. All other **UPDATE STATISTICS** syntax is parsed for compatibility only and ignored. In a batch or stored procedure, only the first **UPDATE STATISTICS** statement for a given table generates a call to **\$SYSTEM.SQL.TuneTable()**. For further details, see [Tune Table](#) in *SQL Optimization Guide*.

If the TSQL [TRACE configuration option](#) is set, the trace log file will contain records of the tables that were tuned.

4.8.7 USE database

Supported, also an extension: **USE NONE** to select no database. Effective at generation-time, persists as long as the transform object exists (e.g. in the shell or loading a batch).

4.9 InterSystems Extensions

TSQL supports a number of InterSystems extensions to Transact-SQL. To allow for the inclusion of these InterSystems-only statements in portable code, InterSystems TSQL also supports a special form of the single-line comment: two hyphens followed by a vertical bar. This operator is parsed as a comment by Transact-SQL implementations, but is parsed as an executable statement in InterSystems TSQL. For further details, refer to the Comments section of the [TSQL Constructs](#) chapter of this document.

TSQL includes the following InterSystems extensions:

4.9.1 OBJECTSCRIPT

This extension allows you to include ObjectScript code or InterSystems SQL code in the compiled output. It takes one or more lines of InterSystems code inside curly brackets.

The following Dynamic SQL example uses **OBJECTSCRIPT** because TSQL does not support the InterSystems SQL **%STARTSWITH** predicate:

```
SET myquery = "OBJECTSCRIPT {SELECT Name FROM Sample.Person WHERE Name %STARTSWITH 'A'}"
SET tStatement = ##class(%SQL.Statement).%New(,,"Sybase")
WRITE "language mode set to ",tStatement.%Dialect,!
SET qStatus = tStatement.%Prepare(myquery)
IF qStatus'=1 {WRITE "%Prepare failed:" DO $System.Status.DisplayError(qStatus) QUIT}
SET rset = tStatement.%Execute()
DO rset.%Display()
WRITE !,"End of data"
```

The following Dynamic SQL example uses **OBJECTSCRIPT** to include ObjectScript code in a TSQL routine:

```
SET newtbl=2
SET newtbl(1)="CREATE TABLE Sample.MyTest(Name VARCHAR(40),Age INTEGER)"
SET newtbl(2)="OBJECTSCRIPT {DO $SYSTEM.SQL.TuneTable("Sample.MyTest") WRITE "TuneTable Done",!}"

SET tStatement = ##class(%SQL.Statement).%New(,,"Sybase")
WRITE "language mode set to ",tStatement.%Dialect,!
SET qStatus = tStatement.%Prepare(.newtbl)
IF qStatus'=1 {WRITE "%Prepare failed:" DO $System.Status.DisplayError(qStatus) QUIT}
SET rset = tStatement.%Execute()
DO rset.%Display()
WRITE !,"End of data"
```

Note that in the above example the **WRITE** command specifies a new line (!); this is necessary because the **OBJECTSCRIPT** extension does not issue a new line following execution.

4.9.2 IMPORTASQUERY

This extension forces a stored procedure to be imported as a query rather than as a class method. This is useful for stored procedures that contain only an EXEC statement, because InterSystems IRIS cannot otherwise determine at import whether such a stored procedure is a query or not.

5

TSQL Settings

Settings are used to tailor the behavior of the compiler and colorizer. The TSQL configuration options are part of the standard InterSystems IRIS configuration.

InterSystems IRIS® data platform supports the following TSQL settings:

- [DIALECT](#)
- [ANSI_NULLS](#)
- [CASEINSCOMPARE](#) (String comparison is not case-sensitive.)
- [QUOTED_IDENTIFIER](#)
- [Equal Literal Replacement](#)
- [TRACE](#)

These values are used to set the corresponding ^%SYS("tsql","SET",...) global array values.

For further details, see [TSQL Compatibility](#) in the *Configuration Parameter File Reference*.

You can view and modify these settings using the InterSystems IRIS Management Portal and/or the %SYSTEM.TSQL Get and Set class methods.

- Go into the InterSystems IRIS Management Portal. Go to **System Administration, Configuration, SQL and Object Settings, TSQL Compatibility**. Here you can specify the **DIALECT** (Sybase or MSSQL, default is Sybase), and turn on or off the **ANSI_NULLS**, **CASEINSCOMPARE**, and **QUOTED_IDENTIFIER** settings. You must press the **Save** button for configuration changes to take effect.
- Invoke the `$$SYSTEM.TSQL.CurrentSettings()` method to display the settings:

```
DO ##class(%SYSTEM.TSQL).CurrentSettings()
```

You can use %SYSTEM.TSQL class methods to get or set the ANSI_NULLS, CaseInsCompare, Quoted_Identifier and EqualLiteralReplacement settings. These methods take a dialect string and change both the current dialect and the specified setting. There are not separate settings for each TSQL dialect. For example, changing CaseInsCompare changes this configuration setting for both Sybase and MSSQL.

5.1 DIALECT

The DIALECT configuration option allows you to select the Transact-SQL dialect. The available options are Sybase and MSSQL. The default is Sybase. You can return the current setting using `$$SYSTEM.TSQL.GetDialect()`. This option is set system-wide using the InterSystems IRIS Management Portal or by using the `$$SYSTEM.TSQL.SetDialect()` method:

```
WRITE ##class(%SYSTEM.TSQL).SetDialect("Sybase")
```

This method returns the prior Dialect setting.

If DIALECT=MSSQL: a DECLARE statement binds host variable values.

If DIALECT=Sybase: host variable values are refreshed for each cursor OPEN.

5.2 ANSI_NULLS

The ANSI_NULLS configuration option allows you to specify whether comparisons to a null value return true or false.

The default is OFF.

- **ON:** All comparisons to a null value evaluate to Unknown. For example, `Age = Null` returns false, even when Age is null. Null is unknown, so it is false/unknown to specify `null=null`.
- **OFF:** Comparisons of a non-Unicode value to a null value evaluates to True if both values are null. For example: `Age = Null` returns true for null values for Age.

You can determine the current ANSI_NULLS setting using `%SYSTEM.TSQL` class methods, or from the *TSQLAnsiNulls* property, as follows:

```
SET context=##class(%SYSTEM.Context.SQL).%New()  
WRITE "ANSI_NULLS is = ",context.TSQLAnsiNulls
```

You can return the current setting using `$$SYSTEM.TSQL.GetAnsiNulls()`. This method returns both the current default dialect and the current ANSI_NULLS setting as a comma-separated string: for example, `MSSQL, ON`.

You can activate (ON) or deactivate (OFF) ANSI_NULLS system-wide using the InterSystems IRIS Management Portal or by using the `$$SYSTEM.TSQL.SetAnsiNulls()` method:

```
WRITE ##class(%SYSTEM.TSQL).SetAnsiNulls("Sybase","OFF")
```

This method sets both the default dialect and the ANSI_NULLS setting, returns the prior settings as a comma-separated string: for example, `MSSQL, ON`.

5.3 CASEINSCOMPARE

The CASEINSCOMPARE setting specifies non-case-sensitive equality comparisons, such as `'A'='a'`. The default is OFF. If this option is set to ON, the comparison operators `=` and `<>` operate without regard to case in most contexts. However, there are a few contexts where such insensitivity does not apply:

- Where a comparison is the ON condition for a JOIN.
- Where either operand is a subquery.

These exceptions exist because InterSystems SQL does not accept the %SQLUPPER operator in these contexts.

You can determine the current CASEINSCOMPARE setting using %SYSTEM.TSQL class methods, or from the *TSQLCaseInsCompare* property, as follows:

```
SET context=##class(%SYSTEM.Context.SQL).%New()
WRITE "ANSI_NULLS is = ",context.TSQLCaseInsCompare
```

You can return the current setting using **\$SYSTEM.TSQL.GetCaseInsCompare()**. You can activate (ON) or deactivate (OFF) CASEINSCOMPARE system-wide using the InterSystems IRIS Management Portal or by using the **\$SYSTEM.TSQL.SetCaseInsCompare()** method:

```
WRITE ##class(%SYSTEM.TSQL).SetCaseInsCompare("Sybase","OFF")
```

This method returns the prior CASEINSCOMPARE setting.

5.4 QUOTED_IDENTIFIER

The QUOTED_IDENTIFIER configuration option allows you to select whether quoted identifiers are supported. The default is OFF (not supported). This option is set using the InterSystems IRIS Management Portal. When QUOTED_IDENTIFIER is on, double quotes are parsed as delimiting an identifier. When QUOTED_IDENTIFIER is off, double quotes are parsed as alternative delimiters for string literals. The preferable delimiters for string literals are single quotes.

You can determine the current QUOTED_IDENTIFIER setting using %SYSTEM.TSQL class methods, or from the *TSQLQuotedIdentifier* property, as follows:

```
SET context=##class(%SYSTEM.Context.SQL).%New()
WRITE "ANSI_NULLS is = ",context.TSQLQuotedIdentifier
```

You can return the current setting using **\$SYSTEM.TSQL.GetQuotedIdentifier()**. You can activate (ON) or deactivate (OFF) QUOTED_IDENTIFIER system-wide using the InterSystems IRIS Management Portal or by using the **\$SYSTEM.TSQL.SetQuotedIdentifier()** method:

```
WRITE ##class(%SYSTEM.TSQL).SetQuotedIdentifier("Sybase","OFF")
```

This method returns the prior QUOTED_IDENTIFIER setting.

5.5 Equal Literal Replacement

The Equal Literal Replacement configuration option is set using a method; is not available from the Management Portal. It controls the behavior of the TSQL compiler. You can return the current setting using **\$SYSTEM.TSQL.GetEqualLiteralReplacement()**. The default is ON.

You can activate (ON) or deactivate (OFF) Equal Literal Replacement system-wide using the **\$SYSTEM.TSQL.SetEqualLiteralReplacement()** method:

```
WRITE ##class(%SYSTEM.TSQL).SetEqualLiteralReplacement("Sybase","OFF")
```

Setting **SetEqualLiteralReplacement("Sybase","OFF")** means TSQL queries with a WHERE clause equal sign (=) predicate or an [IN\(...\) predicate](#) will not perform [literal substitution](#) for literal values on the left or right side of the equal sign, or for any literal value in the IN predicate. This can help the query optimizer choose a better plan when the condition includes fields that have an [outlier value](#).

5.6 TRACE

The TRACE configuration option is not available from the Management Portal. It controls the behavior of the TSQL compiler. When a method is compiled with TRACE active, running this method will log traced messages to the active log file, by default. The active log file is located in the same namespace as IRIS.DAT, and is named using the current process number.

You can activate (1) or deactivate (0) TRACE system-wide using the following ObjectScript command:

```
SET ^%SYS( "tsql" , "TRACE" )=1
```

To return the current trace setting:

```
WRITE ^%SYS( "tsql" , "TRACE" )
```

The TRACE log file records a timestamp for each operation, the elapsed time for each operation, a global references count and a %ROWCOUNT (where applicable). Note that TRUNCATE TABLE always returns a %ROWCOUNT of -1. If an operation involves sharded tables, the global references count is only for the process the procedure is executed on. Work sent to the other shards is not included in the global reference count.

The TRACE log file represents a temporary table using the internal temporary table name; it displays the corresponding user-specified #TempTable name in a /* mytemptable */ comment.

The following is a typical TRACE log file:

Log restarted: 07/26/2019 14:05:02

07/26/2019 14:05:02.324831
CALLSP::testindproc()

07/26/2019 14:05:02.325276
PREPARE TRUNCATE: TRUNCATE TABLE TmpOrder

%ROWCOUNT = -1

Elapsed time = .112758s

07/26/2019 14:05:02.438113
PREPARE INSERT: INSERT INTO TmpOrder (orderid) SELECT OrderId FROM ExOrder WHERE OrderId IN (SELECT DISTINCT OrderId FROM ExFill)

%ROWCOUNT = 956808 Elapsed time = 9.781208s
RETURN:: testindproc with value = 1

07/26/2019 14:05:12.219431

context object: 200@%Library.ProcedureContext
PREPARE DROP TABLE: DROP TABLE IRIS_Temp.LTT75 /* #tab2 */

6

TSQL Functions

6.1 Supported Functions

The following TSQL functions are implemented for InterSystems IRIS® data platform.

6.1.1 ABS

```
ABS(num)
```

Returns the absolute value of *num*. Thus both 123.99 and –123.99 return 123.99.

6.1.2 ACOS

```
ACOS(float)
```

Arc cosine: returns the angle in radians whose cosine is *float*. Thus 1 returns 0.

6.1.3 ASCII

```
ASCII(char)
```

Returns the integer value corresponding to the first character in string *char*. Thus, ASCII ('A') returns 65.

ASCII is functionally identical to **UNICODE**. The reverse of this function is **CHAR**.

6.1.4 ASIN

```
ASIN(float)
```

Arc sine: returns the angle in radians whose sine is *float*. Thus 1 returns 1.570796326...

6.1.5 ATAN

```
ATAN(float)
```

Arc tangent: returns the angle in radians whose tangent is *float*. Thus 1 returns .785398163...

6.1.6 AVG

```
AVG(numfield)
AVG(DISTINCT numfield)
```

Aggregate function: used in a query to return the average of the values in the *numfield* column. For example, `SELECT AVG(Age) FROM Sample.Person`. `AVG(DISTINCT numfield)` averages the number of unique values in the *numfield* column. Fields with NULL are ignored.

6.1.7 CAST

```
CAST(expression AS datatype)
```

Returns the *expression* converted to the specified *datatype*. **CAST** can be used with any supported data type. For further details, refer to [Data Types](#) in the *InterSystems SQL Reference*. **CAST** supports user-defined data types created using the `sp_addtype` stored procedure.

When *expression* is a date value string, such as '2004-11-23' and *datatype* is `TIMESTAMP` or `DATETIME`, a time value of '00:00:00' is supplied.

When *expression* is a time value string, such as '1:35PM' and *datatype* is `TIMESTAMP` or `DATETIME`, the time is converted to a 24-hour clock, the AM or PM suffix is removed, a missing seconds interval is filled in with zeros, and the default date value of '1900-01-01' is supplied. Thus '1:35PM' is converted to '1900-01-01 13:35:00'.

When *expression* is a date value string, such as '2004-11-23' and *datatype* is `DATE`, the date is returned in InterSystems IRIS [\\$HOROLOG](#) date format, such as 60703 (March 14, 2007).

InterSystems TSQL does not support data type XML. However, instead of generating an error during compilation, `CAST(x AS XML)` in SQL mode generates `CAST(x AS VARCHAR(32767))`. In procedure mode, `CAST(x AS XML)` does not generate any conversion.

See **CONVERT**.

6.1.8 CEILING

```
CEILING(num)
```

Returns the closest integer greater than or equal to *num*. Thus 123.99 returns 124, -123.99 returns -123.

The Sybase CEIL synonym is not supported.

6.1.9 CHAR

```
CHAR(num)
```

Returns the character corresponding to the integer value *num*. Thus `CHAR(65)` returns A.

CHAR is functionally identical to **NCHAR**. The reverse of this function is **ASCII**.

6.1.10 CHAR_LENGTH / CHARACTER_LENGTH

```
CHAR_LENGTH(string)
CHARACTER_LENGTH(string)
```

Returns the number of characters in *string*.

6.1.11 CHARINDEX

```
CHARINDEX(seekstring,target [,startpoint])
```

Returns the position in *target* (counting from 1) corresponding to first character of the first occurrence of *seekstring*. You can use the optional *startpoint* integer to specify where to begin the search. The return value counts from the beginning of *target*, regardless of the *startpoint*. If *startpoint* is not specified, specified as 0, 1, or as a negative number, *target* is searched from the beginning. **CHARINDEX** returns 0 if *seekstring* is not found.

6.1.12 COALESCE

```
COALESCE(expression1,expression2,...)
```

Returns the first non-null expression from the specified list of expressions.

6.1.13 COL_NAME

```
COL_NAME(object_id,column_id)
```

Returns the name of the column. Can be used in procedure code or trigger code.

TSQL supports the two-argument form of this function. It does not support a third argument.

The following example returns the column name of the 4th column of Sample.Person:

```
SET sql=2
SET sql(1)="SELECT 'column name'=COL_NAME(id,4) FROM Sample.Person"
SET sql(2)="WHERE id=OBJECT_ID('Sample.Person')"
SET statement=##class(%SQL.Statement).%New()
SET statement.%Dialect="MSSQL"
SET status=statement.%Prepare(.sql)
SET result=statement.%Execute()
DO result.%Display()
```

COL_NAME does not support the Sybase third argument.

6.1.14 CONVERT

```
CONVERT(datatype,expression [,style])
```

Returns the *expression* converted to the specified *datatype*.

- **BIT:** When *datatype* is BIT and *expression* is a boolean value: if the input value is a non-zero number, the result is 1. if the input value is 0, the result is 0. If the input value is the string 'TRUE' (case insensitive), the result is 1. If the input value is the string 'FALSE' (case insensitive), the result is 0. If the input value is NULL, the result is NULL. Any other input value generates an SQLCODE -141 error.
- **CHAR:** When the length of CHAR is not specified, converts to a length of 30 characters.

- DATETIME: When *datatype* is datetime or timestamp:
 - When *expression* is a date value string, such as '2004-11-23', a time value of '00:00:00' is supplied.
 - When *expression* is a time value string, such as '1:35PM' and *datatype* is datetime or timestamp, the time is converted to a 24-hour clock, the AM or PM suffix is removed, a missing seconds interval is filled in with zeros, and the default date value of '1900-01-01' is supplied. Thus '1:35PM' is converted to '1900-01-01 13:35:00'.
 - When *expression* is the empty string or a string of one or more blank spaces, **CONVERT** returns '1900-01-01 00:00:00'.

By default DATETIME data type is mapped to %Library.TimeStamp. **CONVERT** also supports data type mapping of DATETIME to %Library.PosixTime.

CONVERT supports the DATETIME2 data type. InterSystems IRIS maps DATETIME2 to system-defined DDL mapping %Library.TimeStamp. This mapping is supplied with new installs; if you are using an upgrade install, you may need to create this mapping.

CONVERT supports user-defined data types created using the [sp_addtype](#) stored procedure.

The optional *style* argument is used to specify a date/time format when converting a datetime or timestamp value to a string. By specifying various *style* codes you can return a dates and times in a variety of different formats. The available *style* codes are 100 through 116, 120 through 123, 126, 130 and 131, 136 through 140. (the corresponding codes 0 through 7 and 10 through 40 return the same values with two-digit years); The default *style* for a datetime is 0:

```
mon dd yyyy hh:mmAM
```

The following are some of the supported datetime styles for Sybase:

```
15 / 115 = format dd/[yy]yy/mm
16 / 116 = format mon dd yyyy HH:mm:ss
22 / 122 = format [yy]yy/mm/dd HH:mm AM (or PM)
23 / 123 = format [yy]yy-mm-ddTHH:mm:ss
36 / 136 = format hh:mm:ss.zzzzzzAM(PM)
37 / 137 = format hh:mm:ss.zzzzzz
38 / 138 = format mon dd [yy]yy hh:mm:ss.zzzzzzAM(PM)
39 / 139 = format mon dd [yy]yy HH:mm:ss.zzzzzz
40 / 140 = format yyyy-mm-dd hh:mm:ss.zzzzzz
```

The 20 & 21 (120 & 121) *style* codes return the ODBC timestamp format; 20 truncates to whole seconds, 21 returns fractional seconds:

```
yyyy-mm-dd hh:mm:ss.fff
```

For further details, refer to the functionally identical InterSystems SQL [CONVERT](#) function in the *InterSystems SQL Reference*.

See **CAST**.

6.1.15 COS

```
COS(float)
```

Cosine: returns the cosine of the angle specified in *float*. Thus 1 returns .540302305...

6.1.16 COT

```
COT(float)
```

Cotangent: returns the cotangent of the angle specified in *float*. Thus 1 returns .64209261593...

6.1.17 COUNT

```
COUNT(field)
COUNT(DISTINCT field)
COUNT(*)
COUNT(1)
```

Aggregate function: used in a query to return the count of the values in the *field* column. Fields with NULL are not counted. For example, `SELECT COUNT(Name) FROM Sample.Person`. `COUNT(*)` and `COUNT(1)` are synonyms, they count all rows. `COUNT(DISTINCT field)` counts the number of unique values in the *field* column. Fields with NULL are not counted.

6.1.18 CURRENT_DATE

```
CURRENT_DATE
CURRENT DATE
```

Returns the current local date in the following format:

```
yyyy-mm-dd
```

The two syntax forms, with and without an underscore, are identical. Note that no parentheses are used with this function. This function is provided for compatibility with SQL Anywhere; it is supported by both the Sybase and MSSQL dialects.

6.1.19 CURRENT_TIME

```
CURRENT_TIME
CURRENT TIME
```

Returns the current local time in the following format:

```
hh:mm:ss
```

Time is specified using a 24-hour clock, Fractional seconds are not returned.

The two syntax forms, with and without an underscore, are identical. Note that no parentheses are used with this function. This function is provided for compatibility with SQL Anywhere; it is supported by both the Sybase and MSSQL dialects.

6.1.20 CURRENT_TIMESTAMP

```
CURRENT_TIMESTAMP
CURRENT TIMESTAMP
```

Returns the current local date and time in the following format:

```
yyyy-mm-dd hh:mm:ss
```

Time is specified using a 24-hour clock, Fractional seconds are not returned.

The two syntax forms, with and without an underscore, are identical. Note that no parentheses are used with this function.

6.1.21 CURRENT_USER

```
CURRENT_USER
```

Returns the name of the current user.

Note that no parentheses are used with this function.

6.1.22 DATALENGTH

```
DATALENGTH(expression)
```

Returns an integer specifying the number of bytes used to represent *expression*. Thus 'fred' returns 4, and +007.500 returns 3.

6.1.23 DATEADD

```
DATEADD(code, num, date)
```

Returns the value of *date* modified by adding the interval specified in *code* the *num* number of times. The *date* can be a date, time, or date/time string in a variety of formats. You can specify any of the following *code* values, either the abbreviation (left column) or the name (right column):

yy	Year
qq	Quarter
mm	Month
dy	DayofYear
dd	Day
dw, w	Weekday
wk	Week
hh	Hour
mi	Minute
ss	Second
ms	Millisecond

Code values are not case-sensitive. Day, DayofYear, and Weekday all return the same value.

The value returned by DATEADD always includes both date and time in the format:

```
yyyy-mm-dd hh:mm:ss.n
```

Fractional seconds are only returned if the source contained fractional seconds.

If a date is not specified (that is, if *date* contains only a time value), it defaults to 1/1/1900.

If a time is not specified in *date*, it defaults to 00:00:00. Hours are always returned based on a 24-hour clock.

6.1.24 DATEDIFF

`DATEDIFF(code,startdate,enddate)`

Returns the number of *code* intervals between *startdate* and *enddate*. The two dates can be a date, a time, or a date/time string. in the following format:

`yyyy-mm-dd hh:mm:ss.n`

You can specify any of the following *code* values, either the abbreviation (left column) or the name (right column):

yy	Year
mm	Month
dd	Day
dw, w	Weekday
wk	Week
hh	Hour
mi	Minute
ss	Second
ms	Millisecond

Code values are not case-sensitive. Day, DayofYear, and Weekday all return the same value.

If a date is not specified (that is, if *startdate* or *enddate* contains only a time value), it defaults to 1/1/1900.

If a time is not specified in *startdate* or *enddate*, it defaults to 00:00:00.

6.1.25 DATENAME

`DATENAME(code,date)`

Returns the value of the part of the date specified by *code* as a string. The *date* can be a date, time, or date/time string in a variety of formats. *date* must be specified as a quoted string; *code* permits, but does not require enclosing quotes. Available *code* values are:

yyyy, yy year	Year. Returns a four-digit year. If a two-digit year is specified, DATENAME supplies '19' as first two digits.
qq, q quarter	Quarter. Returns an integer 1 through 4.
mm, m month	Month. Returns the full name of the month. For example, 'December'.
dy, y dayofyear	Day of Year. Returns an integer count of days 1 through 366.
dd, d day	Day of Month. Returns an integer count 1 through 31.
wk, ww week	Week of Year. Returns an integer count 1 through 53.
dw, w weekday	Day of Week. Returns the number of the day of the week, counting from Sunday. For example, 3 is Tuesday.
hh hour	Hour. Returns the hour of the day (24-hour clock), an integer 0 through 23.
mi, n minute	Minute. Returns an integer 0 through 59.
ss, s second	Second. Returns a decimal number 0 through 59 which may have a fractional part representing milliseconds.
ms millisecond	Millisecond. Returns the fractional part of a second as an integer.

Code values are not case-sensitive.

If a date is not specified, it defaults to 1/1/1900. Two-digit years default to 19xx.

If a time is not specified, it defaults to 00:00:00. Hours are always returned based on a 24-hour clock. Seconds are always returned with fractional seconds, if fractional seconds are defined. Milliseconds are returned as an integer, not a decimal fraction.

6.1.26 DATEPART

```
DATEPART(code, date)
```

Returns the value of the part of the date specified in *code* as an integer. The *date* can be a date, time, or date/time string in a variety of formats. Available *code* values are listed in **DATENAME**.

6.1.27 DAY

`DAY(date)`

Returns the day portion of the specified date or date/time string. The *date* can be specified in ODBC timestamp format:

`yyyy-mm-dd hh:mm:ss.n`

The *date* must contain a date component. The date separator must be a hyphen (-).

The *date* can also be specified in InterSystems IRIS [\\$HOROLOG](#) date format, such as 60703 (March 14, 2007).

6.1.28 DB_NAME

`DB_NAME ()`

Returns the current namespace name. No argument is permitted.

6.1.29 DEGREES

`DEGREES(float)`

Converts an angle measurement in radians to the corresponding measurement in degrees.

6.1.30 ERROR_MESSAGE

When invoked from within a CATCH block, returns the current error message. Otherwise, returns NULL.

6.1.31 ERROR_NUMBER

When invoked from within a CATCH block, returns the current SQLCODE error. Otherwise, returns NULL.

6.1.32 EXEC

`EXEC(@var)`

Executes dynamic SQL at runtime, as shown in the following example:

```
DECLARE @dyncode VARCHAR(200)
SELECT @dyncode='SELECT TOP 4 Name, Age FROM Sample.Person'
EXEC(@dyncode)
```

Compare this dynamic execution with the [EXECUTE](#) command that executes a stored procedure.

6.1.33 EXP

`EXP(num)`

Returns the exponential of *num*. This is the e constant (2.71828182) raised to the power of *num*. Thus EXP(2) returns 7.3890560989.

6.1.34 FLOOR

```
FLOOR(num)
```

Returns the closest integer less than or equal to *num*. Thus 123.99 returns 123, −123.99 returns −124.

6.1.35 GETDATE

```
GETDATE ( )
```

Returns the current local date and time in the following format:

```
yyyy-mm-dd hh:mm:ss.n
```

Time is specified using a 24-hour clock, Fractional seconds are returned.

6.1.36 GETUTCDATE

```
GETUTCDATE ( )
```

Returns the current UTC (Greenwich Mean Time) date and time in the following format:

```
yyyy-mm-dd hh:mm:ss.n
```

Time is specified using a 24-hour clock, Fractional seconds are returned.

6.1.37 HOST_NAME

```
HOST_NAME ( )
```

Returns the system name of the current host system.

6.1.38 INDEX_COL

```
INDEX_COL(table_name,index_id,key,[,user_id])
```

Returns the name of the indexed column in the specified table. *table_name* can be fully qualified. *index_id* is the number of the table's index. *key* is a key in the index, a value between 1 and sysindexes.keycnt (for a clustered index) or sysindexes.keycnt+1 (for a non-clustered index). *user_id* is parsed but ignored.

6.1.39 ISNULL

```
ISNULL(expr,default)
```

If *expr* is NULL, returns *default*. If *expr* is not NULL, returns *expr*.

6.1.40 ISNUMERIC

```
ISNUMERIC(expression)
```

A boolean function that returns 1 if *expression* is a valid numeric value; otherwise, returns 0.

If the specified *expression* is a field with a null value, ISNUMERIC returns null.

6.1.41 LEFT

```
LEFT(string,int)
```

Returns *int* number of characters from *string*, counting from the left. If *int* is larger than *string*, the full string is returned. See RIGHT.

6.1.42 LEN

```
LEN(string)
```

Returns the number of characters in *string*.

6.1.43 LOG

```
LOG(num)
```

Returns the natural logarithm of *num*. Thus LOG(2) returns .69314718055.

6.1.44 LOG10

```
LOG10(num)
```

Returns the base-10 logarithm of *num*. Thus LOG10(2) returns .301029995663.

6.1.45 LOWER

```
LOWER(string)
```

Returns *string* with all uppercase letters converted to lowercase. See UPPER.

6.1.46 LTRIM

```
LTRIM(string)
```

Removes leading blanks from *string*.

If *string* consists entirely of blank spaces, the dialect determines behavior:

- Sybase: returns NULL.
- MSSQL: returns the empty string.

See [RTRIM](#).

6.1.47 MAX

`MAX(numfield)`

Aggregate function: used in a query to return the largest (maximum) of the values in the *numfield* column. For example:

```
SELECT MAX(Age) FROM Sample.Person
```

Fields with NULL are ignored.

6.1.48 MIN

`MIN(numfield)`

Aggregate function: used in a query to return the smallest (minimum) of the values in the *numfield* column. For example:

```
SELECT MIN(Age) FROM Sample.Person
```

Fields with NULL are ignored.

6.1.49 MONTH

`MONTH(date)`

Returns the month portion of the specified date or date/time string. The *date* can be specified in ODBC timestamp format:

`yyyy-mm-dd hh:mm:ss.n`

The date separator must be a hyphen (-). Dates in any other format return 0.

The *date* can also be specified in InterSystems IRIS [\\$HOROLOG](#) date format, such as 60703 (March 14, 2007).

6.1.50 NCHAR

`NCHAR(num)`

Returns the character corresponding to the integer value *num*. Thus `NCHAR(65)` returns A.

NCHAR is functionally identical to **CHAR**. The reverse of this function is **ASCII**.

6.1.51 NEWID

`NEWID()`

Returns a unique value of a type compatible with the SQL Server UNIQUEIDENTIFIER data type. UNIQUEIDENTIFIER is a system-generated 16-byte binary string, also known as a [globally unique ID \(GUID\)](#). A GUID is used to synchronize databases on occasionally connected systems. A GUID is a 36-character string consisting of 32 hexadecimal numbers separated into five groups by hyphens. InterSystems TSQL does not support UNIQUEIDENTIFIER; it instead uses VARCHAR(36) as the data type for a Globally Unique ID.

The **NEWID** function takes no arguments. Note that the argument parentheses are required.

NEWID() can be used to specify the DEFAULT value when defining a field.

The corresponding InterSystems SQL function is [\\$TSQL_NEWID](#):

```
SELECT $TSQL_NEWID()
```

6.1.52 NOW

```
NOW( * )
```

Returns the current local date and time in the following format:

```
yyyy-mm-dd hh:mm:ss
```

Time is specified using a 24-hour clock, Fractional seconds are not returned.

Note that the asterisk within the parentheses is required.

6.1.53 NULLIF

```
NULLIF(expr1,expr2)
```

Returns NULL if *expr1* is equivalent to *expr2*. Otherwise, returns *expr1*.

6.1.54 OBJECT_ID

```
OBJECT_ID(objname,objtype)
```

Takes the object name as a quoted string, and optionally the object type, and returns the corresponding object ID of the specified object as an integer. The available *objtype* values are as follows: RI = FOREIGN KEY constraint; K = PRIMARY KEY or UNIQUE constraint; P = Stored procedure; S = System table; TR = Trigger; U = User table; V = View.

```
CREATE PROCEDURE GetName
AS SELECT OBJECT_ID('Person','U')
GO
```

Returns the NULL if *objname* does not exist, or if the optional *objtype* is specified and does not match the *objname*. Can be used within procedure code or trigger code. The inverse of OBJECT_NAME.

6.1.55 OBJECT_NAME

```
OBJECT_NAME(id)
```

Takes the object ID integer and returns the corresponding object name of the specified object. Returns the empty string if *id* does not exist. Can be used within procedure code or trigger code. The inverse of OBJECT_ID.

```
CREATE PROCEDURE GetID
AS SELECT OBJECT_NAME(22)
GO
```

6.1.56 PATINDEX

```
PATINDEX(pattern,string)
```

Returns an integer specifying the beginning position of the first occurrence of *pattern* in *string*, counting from 1. If *pattern* is not found in *string*, 0 is returned. Specify *pattern* as a quoted string. Comparisons are case-sensitive. The *pattern* can contain the following wildcard characters:

%	Zero or more characters. For example, '%a%' returns the position of the first occurrence of 'a' in <i>string</i> , including 'a' as the first character in string.
_	Any single character. For example, '_l%' returns 1 if <i>string</i> begins with a substring such as 'Al', 'el', and 'il'.
[xyz]	Any single character from the specified list of characters. For example, '[ai]l%' returns 1 if <i>string</i> begins with the substring 'al' or 'il', but not 'el' or 'Al'.
[a-z]	Any single character from the specified range of characters. For example, '%s[a-z]t%' matches 'sat', 'set', and 'sit'. A range must be specified in ascending ASCII sequence.

The caret (^) character is not a wildcard character; if included within square brackets it is treated as a literal. A *pattern* commonly consists of a search string enclosed in percent (%) characters '%Chicago%' indicating that the entire *string* should be searched.

PATINDEX is supported for sharded, parallel, and linked table queries.

6.1.57 PI

PI ()

Returns the constant pi. The parentheses are required; no argument is permitted. Thus PI() returns 3.141592653589793238.

6.1.58 POWER

POWER (*num*, *exponent*)

Returns the value *num* raised to *exponent*.

6.1.59 QUOTENAME

QUOTENAME (*value*)

Returns *value* as a [delimited identifier](#). TSQL supports double quotes ("value") as delimiter characters. For example:

```
PRINT 123
// returns 123
PRINT QUOTENAME(123)
// returns "123"
```

6.1.60 RADIANS

RADIANS (*float*)

Converts an angle measurement in degrees to the corresponding measurement in radians.

6.1.61 RAND

```
RAND([seed])
```

Returns a random number as a fractional number less than 1. The optional *seed* integer argument is ignored; it is provided for compatibility. If **RAND** is used more than once in a query it returns different random values.

6.1.62 REPLACE

```
REPLACE(target,search,replace)
```

Finds every instance of the *search* string in the *target* string and replaces it with the *replace* string, and returns the resulting string. To remove the *search* string from the *target* string, specify *replace* as an empty string.

6.1.63 REPLICATE

```
REPLICATE(expression,repeat-count)
```

REPLICATE returns a string of *repeat-count* instances of *expression*, concatenated together.

If *expression* is NULL, **REPLICATE** returns NULL. If *expression* is the empty string, **REPLICATE** returns an empty string.

If *repeat-count* is a fractional number, only the integer part is used. If *repeat-count* is 0, **REPLICATE** returns an empty string. If *repeat-count* is a negative number, NULL, or a non-numeric string, **REPLICATE** returns NULL.

6.1.64 REVERSE

```
REVERSE(string)
```

Reverses the order of the characters in *string*.

6.1.65 RIGHT

```
RIGHT(string,int)
```

Returns *int* number of characters from *string*, counting from the right. If *int* is larger than *string*, the full string is returned. See LEFT.

6.1.66 ROUND

```
ROUND(num,length)
```

Returns *num* rounded to the number of decimal digits specified by the integer *length*. If *length* is greater than the number of decimal digits, no rounding is performed. If *length* is 0, *num* is rounded to an integer. If the *length* argument is omitted, it defaults to 0. If *length* is a negative integer, *num* is rounded to the left of the decimal point. A third argument is not accepted by **ROUND**.

6.1.67 RTRIM

`RTRIM(string)`

Removes trailing blanks from *string*.

If *string* consists entirely of blank spaces, the dialect determines behavior:

- Sybase: returns NULL.
- MSSQL: returns the empty string.

See [LTRIM](#).

6.1.68 SCOPE_IDENTITY

Returns the last identity value inserted into an **IDENTITY** column in the same scope. However, the last **IDENTITY** is not limited to the scope of the current procedure. Therefore, you should only use **SCOPE_IDENTITY** when you know that a statement within the current procedure has generated an **IDENTITY** value. For example, **SCOPE_IDENTITY** should be used after an **INSERT** command in the same procedure.

The following Dynamic SQL example returns the **IDENTITY** value from the second **INSERT**:

```
SET sql=6
SET sql(1)="CREATE TABLE #mytest (MyId INT IDENTITY(1,1), "
SET sql(2)="Name VARCHAR(20))"
SET sql(3)="INSERT INTO #mytest(Name) VALUES ('John Smith') "
SET sql(4)="INSERT INTO #mytest(Name) VALUES ('Walter Jones') "
SET sql(5)="PRINT SCOPE_IDENTITY()"
SET sql(6)="DROP TABLE #mytest"
SET statement=##class(%SQL.Statement).%New()
SET statement.%Dialect="MSSQL"
SET status=statement.%Prepare(.sql)
SET result=statement.%Execute()
DO result.%Display()
```

6.1.69 SIGN

`SIGN(num)`

Returns a value indicating the sign of *num*. If *num* is negative (for example, -32), it returns -1. If *num* is positive (for example, 32 or +32), it returns 1. If *num* is zero (for example, 0 or -0), it returns 0.

6.1.70 SIN

`SIN(float)`

Sine: returns the sine of the angle specified in *float*. Thus 1 returns .841470984807...

6.1.71 SPACE

`SPACE(num)`

Returns a string of blank spaces of length *num*.

6.1.72 SQRT

```
SQRT(num)
```

Returns the square root of *num*. Thus SQRT(9) returns 3.

6.1.73 SQUARE

```
SQUARE(num)
```

Returns the square of *num*. Thus SQUARE(9) returns 81.

6.1.74 STR

```
STR(num, [length [, precision]])
```

Returns a string of *length* characters. If the integer *length* is equal to or greater than the number of characters in the numeric *num* (including decimal point and sign characters), **STR** returns *num* converted to a string and padded with leading blanks to make the resulting string of *length* characters.

If the optional integer *precision* is specified, *num* is truncated to the specified number of decimal digits before string conversion. If *precision* is omitted, *num* is truncated to its integer portion. If *precision* is larger than the number of decimal digits, *num* is padded with trailing zeros before string conversion.

If *length* is omitted, it defaults to 10. If *length* is less than the number of characters in *num* (after adjustment by *precision*) a dummy string consisting of all asterisks of *length* number of characters is returned.

6.1.75 STUFF

```
STUFF(string, start, length, replace)
```

Returns *string* with *length* number of characters removed and the *replace* string inserted. The point of removal and insertion is specified by the *start* integer, counting from the beginning of *string*. If *length* is 0, no characters are removed. If *replace* is the empty string, no characters are inserted.

If *start* is greater than the number of characters in *string*, no value is returned. If *start* is 1, *length* number of characters are removed from the beginning of *string* and the *replace* string inserted. If *start* is 0, *length* minus 1 number of characters are removed from the beginning of *string* and the *replace* string inserted.

If *length* is greater than or equal to the number of characters in *string*, the *replace* string is returned. The *replace* string length is not limited by the length of *string* or *length*.

6.1.76 SUBSTRING

```
SUBSTRING(string, start, length)
```

Returns a substring of *string* beginning at the location *start* for the *length* number of characters. If *start* is greater than the length of *string*, or if *length* is 0, no string is returned.

6.1.77 SUM

```
SUM(numfield)  
SUM(DISTINCT numfield)
```

Aggregate function: used in a query to return the sum of the values in the *numfield* column. For example:

```
SELECT SUM(Age) FROM Sample.Person
```

SUM(DISTINCT *numfield*) sums the unique values in the *numfield* column. Fields with NULL are ignored.

6.1.78 SUSER_NAME

```
SUSER_NAME ( )
```

Returns the name of the current OS user. Parentheses are required, no argument is permitted. Equivalent to TSQL **USER_NAME()**, the InterSystems SQL **USER** function, and the ObjectScript **\$USERNAME** special variable.

6.1.79 SUSER_SNAME

```
SUSER_SNAME ( )
```

Returns the name of the current OS user. Parentheses are required, no argument is permitted. Equivalent to TSQL **USER_NAME()**, the InterSystems SQL **USER** function, and the ObjectScript **\$USERNAME** special variable.

6.1.80 TAN

```
TAN(float)
```

Tangent: returns the tangent of the angle specified in *float*. Thus 1 returns 1.55740772465...

6.1.81 TEXTPTR

```
TEXTPTR(field)
```

Returns an internal pointer to the image or text column data specified in *field*. The data type of this pointer is VARBINARY(16).

6.1.82 TEXTVALID

```
TEXTVALID(' table.field' , textpointer)
```

Takes an internal pointer to an image or text column from **TEXTPTR**, and compares it to a specified in *table.field*. Returns 1 if the pointer points to the specified *table.field*. Otherwise, returns 0.

6.1.83 UNICODE

`UNICODE(char)`

Returns the Unicode integer value corresponding to the first character in the string *char*. Thus, `UNICODE('A')` returns 65.

UNICODE is functionally identical to **ASCII**. The reverse of this function is **CHAR**.

6.1.84 UPPER

`UPPER(string)`

Returns *string* with all lowercase letters converted to uppercase. See **LOWER**.

6.1.85 USER

`USER`

Returns the name of the current user.

Note that no parentheses are used with this function.

6.1.86 USER_NAME

`USER_NAME([userid])`

Returns the name of the user specified by user ID. If the optional *userid* is omitted, returns the name of the current user. The argument is optional; the parentheses are mandatory.

6.1.87 YEAR

`YEAR(date)`

Returns the year portion of the specified date or date/time string. The *date* can be specified in ODBC timestamp format:

`yyyy-mm-dd hh:mm:ss.n`

The date separator can be either a hyphen (-) or a slash (/).

The *date* can also be specified in InterSystems IRIS [\\$HOROLOG](#) date format, such as 60703 (March 14, 2007).

6.2 Unsupported Functions

The following Microsoft Transact-SQL functions are not supported by TSQL at this time: APP_NAME, ATN2, BINARY_CHECKSUM, CHECKSUM, COL_LENGTH, COLLATIONPROPERTY, COLUMNPROPERTY, CURSOR_STATUS, DATABASEPROPERTY, DATABASEPROPERTYEX, DB_ID, DIFFERENCE, FILE_ID, FILE_NAME, FILEGROUP_ID, FILEGROUP_NAME, FILEGROUPPROPERTY, FILEPROPERTY, FORMATMESSAGE, FULLTEXTCATALOGPROPERTY, FULLTEXTSERVICEPROPERTY, GETANSINULL, HOST_ID, IDENT_CURRENT, IDENT_INCR, IDENT_SEED, IDENTITY, INDEXKEY_PROPERTY, INDEXPROPERTY, ISDATE, IS_MEMBER,

IS_SRVROLEMEMBER, OBJECTPROPERTY, PARSENAME, PERMISSIONS, ROWCOUNT_BIG, SERVERPROPERTY, SESSIONPROPERTY, SESSION_USER, SOUNDEX, SQL_VARIANT_PROPERTY, STATS_DATE, STDEV, STDEVP, SYSTEM_USER, TYPEPROPERTY.

7

TSQL Variables

7.1 Local Variables

By default, TSQL local variables are specified using an At Sign (@) prefix. For example, @myvar. You can override this default to also allow PLAINLOCALS, TSQL local variables specified without an At Sign (@) prefix. For example, myvar.

7.1.1 Declaring a Local Variable

A local variable must be declared (using **DECLARE** or as a formal parameter) before use. A variable name must be a valid [identifier](#). Local variable names are not case-sensitive. The declaration must specify a data type, though strict data typing is not enforced in InterSystems TSQL. For a list of supported data types, refer to the [TSQL Constructs](#) chapter of this document.

The **DECLARE** command has the following syntax:

```
DECLARE @var [AS] datatype [ = initval]
```

If declaring variables is inconvenient, you can switch this check off using the NDC setting. However, cursors must be declared, even if NDC is used.

Stored procedure arguments are automatically declared as local variables.

7.1.2 Setting a Local Variable

A local variable can be set using either the **SET** command or the **SELECT** command. A local variable can be displayed using either the **PRINT** command or the **SELECT** command. The following Dynamic SQL examples show two local variables being declared, set, and displayed:

```
SET myquery = 3
SET myquery(1) = "DECLARE @a CHAR(20),@b CHAR(20) "
SET myquery(2) = "SET @a='hello ' SET @b='world!' "
SET myquery(3) = "PRINT @a,@b"
SET tStatement = ##class(%SQL.Statement).%New(,,"MSSQL")
SET qStatus = tStatement.%Prepare(.myquery)
IF qStatus'=1 { WRITE "%Prepare failed",$System.Status.DisplayError(qStatus) QUIT}
SET rset = tStatement.%Execute()
DO rset.%Display()
```

```
SET myquery = 3
SET myquery(1) = "DECLARE @a CHAR(20),@b CHAR(20) "
SET myquery(2) = "SELECT @a='hello ', @b='world!'"
SET myquery(3) = "SELECT @a,@b"
SET tStatement = ##class(%SQL.Statement).%New(,,"MSSQL")
SET qStatus = tStatement.%Prepare(.myquery)
IF qStatus'=1 { WRITE "%Prepare failed",$System.Status.DisplayError(qStatus) QUIT}
SET rset = tStatement.%Execute()
DO rset.%Display()
```

7.1.3 Initial and Default Values

By default, **DECLARE** initializes local variables to " " (SQL NULL). Optionally, you can specify an initial value (*initval*) for a local variable in the **DECLARE** command.

If a declared variable is set to the results of a scalar subquery, and the subquery returns no rows, InterSystems TSQL sets the variable to " " (SQL NULL). This default is compatible with MS SQLServer; it is not compatible with Sybase.

7.1.4 Plain Local Variables

By default, local variables require an @ prefix. However, you can specify plain locals, local variables that do not require an @ prefix. The following command activates plain local variables:

```
SET PLAINLOCALS ON
```

You must activate plain local variables before declaring these variables. With plain local variables activated you can declare both local variables with an @ prefix and local variables without an @ prefix. However, you cannot declare two variables that only differ by the @ prefix. For example, @myvar and myvar are considered the same variable. When declaring, selecting, or printing a plain local variable, you can specify the same variable with or without the @ prefix.

Plain local variables follow all of the other TSQL variable conventions.

The following TSQL class method specifies PLAINLOCALS ON and declares and uses both an @ local variable and a plain local variable:

```
ClassMethod Hello() As %String [Language=tsql,ReturnResultsets,SqlProc
]
{ SET PLAINLOCALS ON;
  DECLARE @a CHAR(20),b CHAR(20);
  SET @a='hello ' SET b='world!';
  PRINT @a,b;
}
```

7.2 @@ Special Variables

TSQL special variables are identified by an @@ prefix. @@ variables are system-defined; they cannot be created or modified by user processes. @@ variables are global in scope (available to all processes). They are thus sometimes referred to elsewhere in the Transact-SQL literature as “global variables.” Because the term “global variable” is used widely in InterSystems IRIS and differs significantly in meaning, these TSQL @@ variables are referred to here as “special variables” to avoid confusion.

The following special variables are implemented. Invoking an unimplemented special variable generates a #5001 ' @nnn ' unresolved symbol error or a #5002 <UNDEFINED> error. The corresponding ObjectScript and InterSystems SQL generated code for each special variable is provided:

7.2.1 @@ERROR

Contains the error number of the most recent TSQL error. 0 indicates that no error has occurred. A 0 value is returned when either SQLCODE=0 (successful completion) or SQLCODE=100 (no data, or no more data). To differentiate these two results, use @@SQLSTATUS.

ObjectScript SQLCODE

SQL : SQLCODE

7.2.2 @@FETCH_STATUS

Contains an integer specifying the status of the last FETCH cursor statement. The available options are: 0=row successfully fetched; -1=no data could be fetched; -2 row fetched is missing or some other error occurred. A value of -1 can indicate that there is no data to FETCH, or that the fetch has reached the end of the data.

```
SET myquery = "SELECT @@FETCH_STATUS AS FetchStat"
SET tStatement = ##class(%SQL.Statement).%New(,,"MSSQL")
SET qStatus = tStatement.%Prepare(myquery)
IF qStatus'=1 { WRITE "%Prepare failed",$System.Status.DisplayError(qStatus) QUIT}
SET rset = tStatement.%Execute()
DO rset.%Display()
```

The corresponding InterSystems SQL function is:

```
SELECT $TSQL_FETCH_STATUS()
```

ObjectScript \$Case(\$Get(SQLCODE,0),0:0,100:-1,:-2)

SQL CASE :SQLCODE WHEN 0 THEN 0 WHEN 100 THEN -1 ELSE -2 END

7.2.3 @@IDENTITY

Contains the IDENTITY field value of the most recently inserted, updated, or deleted row.

ObjectScript %ROWID

SQL :%ROWID

7.2.4 @@LOCK_TIMEOUT

Contains an integer specifying the timeout value for locks, in seconds. Lock timeout is used when a resource needs to be exclusively locked for inserts, updates, deletes, and selects. The default is 10.

```
SET myquery = "SELECT @@LOCK_TIMEOUT AS LockTime"
SET tStatement = ##class(%SQL.Statement).%New(,,"MSSQL")
SET qStatus = tStatement.%Prepare(myquery)
IF qStatus'=1 { WRITE "%Prepare failed",$System.Status.DisplayError(qStatus) QUIT}
SET rset = tStatement.%Execute()
DO rset.%Display()
```

The corresponding InterSystems SQL function is:

```
SELECT $TSQL_LOCK_TIMEOUT()
```

ObjectScript LOCK command

SQL SET OPTION LOCK_TIMEOUT

7.2.5 @@NESTLEVEL

Contains an integer specifying the nesting level of the current process.

```
SET myquery = "PRINT @@NESTLEVEL"
SET tStatement = ##class(%SQL.Statement).%New(,,"MSSQL")
SET qStatus = tStatement.%Prepare(myquery)
IF qStatus'=1 { WRITE "%Prepare failed",$System.Status.DisplayError(qStatus) QUIT}
SET rset = tStatement.%Execute()
DO rset.%Display()
```

The corresponding InterSystems SQL function is:

```
SELECT $TSQL_NESTLEVEL()
```

ObjectScript \$STACK

7.2.6 @@ROWCOUNT

Contains the number of rows affected by the most recent **SELECT**, **INSERT**, **UPDATE**, or **DELETE** command. A single-row **SELECT** always returns a @@ROWCOUNT value of either 0 (no row selected) or 1.

When invoking an AFTER statement level trigger, the @@ROWCOUNT value upon entering the trigger is the @@ROWCOUNT immediately prior to the trigger. Rows affected within the scope of the trigger code are reflected in the @@ROWCOUNT value. Upon completion of the trigger code, @@ROWCOUNT reverts to the value immediately prior to the trigger invocation.

ObjectScript %ROWCOUNT

SQL :%ROWCOUNT

7.2.7 @@SERVERNAME

Contains the InterSystems IRIS instance name.

```
SET myquery = "SELECT @@SERVERNAME AS CacheInstance"
SET tStatement = ##class(%SQL.Statement).%New(,,"MSSQL")
SET qStatus = tStatement.%Prepare(myquery)
IF qStatus'=1 { WRITE "%Prepare failed",$System.Status.DisplayError(qStatus) QUIT}
SET rset = tStatement.%Execute()
DO rset.%Display()
```

The corresponding InterSystems SQL function is:

```
SELECT $TSQL_SERVERNAME()
```

ObjectScript \$PIECE(\$system,":",2)

7.2.8 @@SPID

Contains the server process ID of the current process.

```
SET myquery = "SELECT @@SPID AS ProcessID"
SET tStatement = ##class(%SQL.Statement).%New(,,"MSSQL")
SET qStatus = tStatement.%Prepare(myquery)
IF qStatus'=1 { WRITE "%Prepare failed",$System.Status.DisplayError(qStatus) QUIT}
SET rset = tStatement.%Execute()
DO rset.%Display()
```

The corresponding InterSystems SQL function is:

```
SELECT $TSQL_SPID()
```

ObjectScript \$JOB

7.2.9 @@SQLSTATUS

Contains an integer specifying the completion status of the most recent SQL statement. Available values are: 0=successful completion; 1=failure; 2=no (more) data available.

```
SET myquery = "SELECT @@SQLSTATUS AS SQLStatus"
SET tStatement = ##class(%SQL.Statement).%New(, , "MSSQL")
SET qStatus = tStatement.%Prepare(myquery)
IF qStatus'=1 { WRITE "%Prepare failed",$System.Status.DisplayError(qStatus) QUIT}
SET rset = tStatement.%Execute()
DO rset.%Display()
```

The corresponding InterSystems SQL function is:

```
SELECT $TSQL_SQLSTATUS()
```

ObjectScript \$Case(\$Get(SQLCODE,0),0:0,100:2,:1)

SQL CASE :SQLCODE WHEN 0 THEN 0 WHEN 100 THEN 2 ELSE 1 END

7.2.10 @@TRANCOUNT

Contains the number of currently active transactions.

```
SET myquery = "SELECT @@TRANCOUNT AS ActiveTransactions"
SET tStatement = ##class(%SQL.Statement).%New(, , "MSSQL")
SET qStatus = tStatement.%Prepare(myquery)
IF qStatus'=1 { WRITE "%Prepare failed",$System.Status.DisplayError(qStatus) QUIT}
SET rset = tStatement.%Execute()
DO rset.%Display()
```

The corresponding InterSystems SQL function is:

```
SELECT $TSQL_TRANCOUNT()
```

ObjectScript \$TLEVEL

7.2.11 @@VERSION

Contains the InterSystems IRIS version number and date and time of its installation.

```
SET myquery = "SELECT @@VERSION AS CacheVersion"
SET tStatement = ##class(%SQL.Statement).%New(, , "MSSQL")
SET qStatus = tStatement.%Prepare(myquery)
IF qStatus'=1 { WRITE "%Prepare failed",$System.Status.DisplayError(qStatus) QUIT}
SET rset = tStatement.%Execute()
DO rset.%Display()
```

The corresponding InterSystems SQL function is:

```
SELECT $TSQL_VERSION()
```

ObjectScript \$ZVERSION

8

TSQL System Stored Procedures

InterSystems IRIS provides TSQL system stored procedures to help you administer and track database objects in your system. TSQL system stored procedures can be executed in any namespace and schema. The scope of all stored procedures is the current namespace.

Stored procedures are executed using the TSQL **EXECUTE** or **EXEC** command.

InterSystems TSQL supports the following system stored procedures:

- [sp_addtype](#)
- [sp_droptype](#)
- [sp_procxmode](#)

8.1 sp_addtype

This system stored procedure adds a user-defined data type.

```
sp_addtype typename, phystype [(length) | (precision [, scale])]
[, "identity" | nulltype]
```

typename is a user-defined data type name. *phystype* is the physical datatype on which to base the user-defined datatype. You can optionally specify either **IDENTITY** or *nulltype*. The optional **IDENTITY** keyword specifies that the user-defined data type has the **IDENTITY** property, which by definition is **NOT NULL**. The optional *nulltype* specifies whether a column with this data type should have a **NOT NULL** constraint: options are **NULL** (allow nulls) and **NOT NULL** or **NONULL** (do not allow nulls).

A *typename* can only be used for table columns. A *typename* is defined within the current namespace; mapping of a *typename* across namespaces is not supported.

The SQL compiler looks up type mapping in the DDL Datatype Mapping definitions in the system configuration before looking up in the *typename* table. Therefore, if you define a *typename* using **sp_addtype** named **My_Type**, but you also have a **My_Type** defined in the User DDL Datatype Mappings, the mapping of **My_Type** comes from the User DDL Datatype Mapping definition.

This lookup of *typename* is done at runtime, so *typename* does not need to be defined at DDL compile time.

The following example creates the data type **shortstr**, which requires a value (**NOT NULL**) and this value must be six or less characters:

```
EXEC sp_addtype 'shortstr','varchar(6)','not null'
```

8.2 sp_droptype

This system stored procedure deletes a user-defined data type.

```
sp_droptype typename
```

typename is a user-defined data type name defined within the current namespace.

See [sp_addtype](#).

8.3 sp_procxmode (Sybase only)

This system stored procedure displays or changes the execution mode associated with stored procedures.

```
sp_procxmode [procname [, tranmode]]
```

procname is the name of a stored procedure. *tranmode* is the transaction execution mode: Values are "chained", "unchained" (the default), and "anymode".

When called with no arguments, **sp_procxmode** returns a result set of all procedures defined for the namespace, their user name and transaction mode (1=unchained). When called with just the *procname* argument, **sp_procxmode** returns the user name and transaction mode of the specified procedure.

The following *tranmode* values are supported:

- Chained: implicitly begins a transaction before any data-retrieval or modification statement: delete, insert, open, fetch, select, or update. You must still explicitly end the transaction with commit transaction or rollback transaction. A procedure defined to run in chained mode will set autocommit_off at the beginning of the procedure, and restore the prior setting at the end of the procedure.
- Unchained: (the default) requires the user to explicit begin a transaction before any data-retrieval or modification statement. You must explicitly end the transaction with commit transaction or rollback transaction. A procedure defined to run in unchained mode will set autocommit_on at the beginning of the procedure, and restore the prior setting at the end of the procedure.
- Anymode: If the mode is not defined, or if it is defined as anymode, no change is made to the autocommit setting.

InterSystems IRIS does not support a SET [UN]CHAINED option to change process settings. The setting used is the current setting of the process's auto-commit mode.

InterSystems IRIS does not report an error if a procedure defined as chained is called by a process in autocommit_on mode, nor does it report an error if a procedure defined as unchained is called by a process in autocommit_off mode.

The *tranmode* metadata is not part of the actual method definition. This means that changing the *tranmode* does not require a recompile. It also means that when exporting/importing classes containing TSQL (Sybase) stored procedures, the *tranmode* setting for the procedure is not exported with the class definition. Upon import, if a procedure needs to be defined in Chained mode, you must call `EXEC sp_procxmode 'procname', 'chained'` for the procedure.